University of Warsaw

Faculty of Mathematics, Informatics and Mechanics

Paweł Balawender

Student no. 429141

Practical programming languages capturing complexity classes

Master's thesis
in COMPUTER SCIENCE

Supervisor:
dr hab. Paweł Parys, prof. UW
Institute of Informatics, University of Warsaw

Warsaw, July 2025

Abstract

In this work, I study which features make sense to add to a programming language from the computational complexity perspective. Specifically, I focus on how these features can be used to capture various complexity classes, such as LOGSPACE, PTIME, and PSPACE. By analyzing the expressiveness and limitations of these features, I aim to provide insights into the design of programming languages that are both practical and theoretically sound.

Keywords

Programming languages, Descriptive complexity, Bounded arithmetic, Lean 4, Rocq, Coq, LOGSPACE, PTIME

Thesis domain (Socrates-Erasmus subject area codes)

11.3 Computer Science

Subject classification

F. Theory of computation
F.3. Logics and meanings of programs
F.3.3. Studies of program constructs

Tytuł pracy w języku polskim

Praktyczne języki programowania wyrażające klasy złożoności

Contents

List of Listings

Chapter 1

Introduction

The purpose of this thesis is to find a convenient way of certifying that a given program is in some complexity class. We want to do it by defining such a programming language, that by definition every program written in it has a given complexity.

(i) first, we check if imperative programming is best for capturing complexity classes in Chapter 3. (E.g. think about language with only for i=1..n loops, i.e. primitive recursion) It turns out it is not; we look for better paradigms (computation models) for our problem;

(ii) then, we check why can't we actually just certify complexity by proving in Coq that a given C++ program operates in time $\mathcal{O}(n^2)$: Chapter 4;

(iii) in Chapter 5 we study logics (which are also a form of syntax) that characterize decisional complexity classes purely in terms of how complex is the specification of the problem. This gives us some characterizations of decisive complexity classes, and we can also sometimes define functions by asking "is k-th bit of output f(x) 0 or 1?" repeatedly;

(iv) the purpose of Chapter 6 is to check decisive complexity classes are enouguh for us, or do we need to really distinguish e.g. NP and FNP. Turns out, sometimes we have to! But: we can obtain e.g. FL from L and NC1 reductions. This leads us to a new paradigm, that I named my Github repo after (on 16 November 2023!): having a simpler languauge for reductions, and using an oracle for e.g. P-complete problem. We can get FL, FP and most important circuit complexity classes like functional AC0 etc. this way. But obtaining a language for these NC1 reductions is not that easy neither! We leave the problem of circuit reductions/uniformity to appendix. We proceed to try find FL characterization right away in Linear Logic (Ugo Dal Lago IntML) and in Recursion Theory (Neergaard functional language);

(v) in Chapter 7, we study simple and cool function algebras that capture FL, and FP. Neergaard implemented a programming language for FL. I obtained his code from 20 years ago (he sent me zip on Linkedin), and refreshed it. The author in one of his works declares that it doesn't look possible to add 'pair' type to this language;

(vi) in Chapter 8, we study Ugo Dal Lago's linear-types programming language that captures FL and FNL. Briefly! Because introducing linear types is too much overhead;

(vii) in Chapter 9, we study my contribution! I noticed that it is feasible to extract computational content from proofs conducted in bounded arithmetic (file Extraction.lean in my repo[1]) And you get guarantees on the computational complexity of code exported

---

[1] https://github.com/ruplet/formalization-of-bounded-arithmetic

in this way. So, if you prove a theorem that a function is total (forall x, we can find y) in bounded arithmetic, (and you didnt use proof by contradiction too heavily), you can convert it to a program with certified complexity. Then, the most difficult part - I analyzed multiple tools([2]) and found a way to define bounded arithmetic in lean in such a way that it inter-operates with mathlib + is "believable" (not too much my own machinery). I presented at aitp, and had a visit at inria where i tested their tool designed to formalize proof-theory in Rocq (difficult to use due, not enough engineering put in their system).

---

[2]`https://github.com/ruplet/formalization-of-bounded-arithmetic/blob/main/presentation-seminar-sliwowica/seminar-presentation.pdf`

Chapter 2

Preliminaries

In this chapter we introduce the standard definitions for logic and complexity theory, necessary to fix the language we will use in the rest of the work. We restrict attention to the logical frameworks needed later (single-sorted first-order logic and its two-sorted counterpart) and omit the analogous presentations for propositional calculus or full second-order logic. The semantics is entirely classical (as in: not intuitionistic).

In Section 2.4 we first introduce the so-called *non-uniform* circuit families (e.g. $\mathsf{FAC}^i_{/poly}$), then define a weak notion of uniformity and the appropriate *uniform* circuit families (e.g. FAC^i). The notion of uniformity used in this thesis (cf. Definition 2.4.3) and the results using it are usually not interesting. However, in this work we will not focus too much on that problem. We leave out the complicated divagations about the appropriate (i.e. much weaker) notions of uniformity to Appendix A.

Remark 2.1 (Bibliography). The definitions in Section 2.1 are in the style of [CN06, Section 2B; CN10, Section II.2]. The definitions in Section 2.2 are in the style of [CN06, Section 4B; CN10, Section IV.2]. We will mostly need them for Chapter 9, but also for Chapter 5 — the results discussed in the latter are mostly from [Imm99], where different style of definitions is used. However, [Imm99] doesn't introduce two-sorted logic. The single-sorted definitions differences are mostly negligible and the consistent treatment of single- and two-sorted logic, is important for Chapter 9.

Much effort has been put into considering different versions of the definitions in Section 2.3 to keep them consistent with the definitions of functional complexity classes studied later in Section 6.2, e.g. the class FNP defined in Definition 6.2.3. These are much less standard, and often have no clear consensus in the literature.

The definitions of decisional circuit classes are from [Vol99].

2.1. Single-sorted first-order logic

Definition 2.1.1 (First-order vocabulary and syntax)**.** A *first-order vocabulary* (or *language*) $\mathcal{L}$ consists of:

   (i) for each $n \geq 0$, a (possibly empty) set of n-ary *function symbols*;

   (ii) for each $n \geq 0$, a set of n-ary *predicate symbols*, which is nonempty for at least one n.

We use $f, g, h, \ldots$ as meta-variables for function symbols and $P, Q, R, \ldots$ for predicate symbols. A 0-ary function symbol is called a *constant symbol*. In addition, the following logical symbols are available to build first-order terms and formulas:

(i) an infinite set of *variables*. We use $x, y, z, \ldots$ and sometimes $a, b, c, \ldots$ as meta-variables for variables;

(ii) the connectives $\neg$, $\wedge$, $\vee$ (not, and, or) and logical constants $\bot$, $\top$ (false, true);

(iii) the quantifiers $\forall$, $\exists$ (for all, there exists);

(iv) parentheses (,).

Remark 2.2. Note that we don't introduce the equality sign $=$ as a special symbol in the vocabulary. We will treat equality as a standard binary relation and only require it to be treated specially when reasoning about semantics.

Definition 2.1.2 ($\mathcal{L}$-terms)**.** Let $\mathcal{L}$ be a first-order vocabulary. The set of $\mathcal{L}$-*terms* is defined inductively as follows:

(i) every variable is an $\mathcal{L}$-term;

(ii) if f is an n-ary function symbol of $\mathcal{L}$ and $t_1, \ldots, t_n$ are $\mathcal{L}$-terms, then

$$f(t_1, \ldots, t_n)$$

is an $\mathcal{L}$-term.

Definition 2.1.3 ($\mathcal{L}$-formulas)**.** Let $\mathcal{L}$ be a first-order vocabulary. The set of *first-order formulas in* $\mathcal{L}$ (or $\mathcal{L}$-*formulas*) is defined inductively as follows:

(i) the logical constants $\bot$ and $\top$ are atomic formulas;

(ii) if P is an n-ary predicate symbol in $\mathcal{L}$ and $t_1, \ldots, t_n$ are $\mathcal{L}$-terms, then

$$P(t_1, \ldots, t_n)$$

is an *atomic* $\mathcal{L}$-formula;

(iii) if φ and ψ are $\mathcal{L}$-formulas, then $\neg\varphi$, $(\varphi \wedge \psi)$, and $(\varphi \vee \psi)$ are $\mathcal{L}$-formulas;

(iv) if φ is an $\mathcal{L}$-formula and x is a variable, then $\forall x. \varphi$ and $\exists x. \varphi$ are $\mathcal{L}$-formulas.

For example,

$$(\neg \forall x. P(x) \vee \exists x. \neg P(x)) \quad \text{and} \quad (\forall x. \neg Q(x, f(y)) \wedge \neg \forall z. Q(f(y), z))$$

are $\mathcal{L}$-formulas (for suitable choices of P, Q, and f in $\mathcal{L}$).

Definition 2.1.4 (The language of arithmetic)**.** The *language of arithmetic* is

$$L_A = [\, 0, \, 1, \, +, \, \cdot \, ; \, =, \, \leq \,],$$

where 0 and 1 are constant symbols, $+$ and $\cdot$ are binary function symbols, and $=$ and $\leq$ are binary predicate symbols. We will write these symbols in infix form.

Definition 2.1.5 (Free and bound variables)**.** Let φ be a formula and x a variable. An occurrence of x in φ is *bound* if it lies within a subformula of φ of the form $\forall x. \psi$ or $\exists x. \psi$. Any other occurrence of x in φ is called *free*.

Definition 2.1.6 (Closed terms, closed formulas, sentences)**.** A formula is *closed* if it contains no free occurrence of any variable. A term is *closed* if it contains no variables at all. A closed formula is also called a *sentence*.

Definition 2.1.7 ($\mathcal{L}$-structure)**.** Let $\mathcal{L}$ be a first-order vocabulary. An $\mathcal{L}$-*structure* $\mathcal{M}$ consists of:

(i) a nonempty set M, called the *universe* (variables are intended to range over M);

(ii) for each n-ary function symbol f in $\mathcal{L}$, an associated function $f^{\mathcal{M}} : M^n \to M$;

(iii) for each n-ary predicate symbol P in $\mathcal{L}$, an associated relation $P^{\mathcal{M}} \subseteq M^n$.

Remark 2.3. Note that to "syntactical" relations, we assign "real" relations defined on the underlying elements of the structure. We will want to treat some of these relations specially, e.g. to make sure that the "$=$" relation is interpreted as the actual equality, or that a designated "$\text{SUCC}(x, y)$" relation holds only if the underlying objects are actual natural numbers, for which we have $x + 1 = y$; see e.g. Definition 5.0.4.

Definition 2.1.8 (Object Assignment)**.** Let $\mathcal{M}$ be a structure with universe M. An *object assignment* σ for $\mathcal{M}$ is a mapping from variables to the universe M.

Notation 1. Let x be a variable and $m \in M$. We write $\sigma(m/x)$ for the assignment that is the same as σ except that it maps x to m.

Definition 2.1.9 (Basic Semantic Definition)**.** Let $\mathcal{L}$ be a first-order vocabulary, let $\mathcal{M}$ be an $\mathcal{L}$-structure with universe M, and let σ be an object assignment for $\mathcal{M}$.

Interpretation of terms. Each $\mathcal{L}$-term t is assigned an element $t^{\mathcal{M}}[\sigma] \in M$, defined by structural induction on t:

(i) for each variable x, $x^{\mathcal{M}}[\sigma] = \sigma(x)$;

(ii) $(ft_1 \ldots t_n)^{\mathcal{M}}[\sigma] = f^{\mathcal{M}}(t_1^{\mathcal{M}}[\sigma], \ldots, t_n^{\mathcal{M}}[\sigma])$.

Satisfaction of formulas. For an $\mathcal{L}$-formula φ, the relation

$$\mathcal{M} \models \varphi[\sigma]$$

(read: "$\mathcal{M}$ satisfies φ under σ") is defined by structural induction on φ:

(i) the structure $\mathcal{M}$ satisfies $\top[\sigma]$ and $\mathcal{M} \not\models \bot[\sigma]$;

(ii) for an atomic formula $Pt_1 \ldots t_n$ (with P an n-ary predicate symbol),

$$\mathcal{M} \models (Pt_1 \ldots t_n)[\sigma] \text{ iff } \langle t_1^{\mathcal{M}}[\sigma], \ldots, t_n^{\mathcal{M}}[\sigma] \rangle \in P^{\mathcal{M}};$$

(iii) if $\mathcal{L}$ contains $=$, then for terms s, t,

$$\mathcal{M} \models (s = t)[\sigma] \text{ iff } s^{\mathcal{M}}[\sigma] = t^{\mathcal{M}}[\sigma];$$

(iv) the structure satisfies $\neg\varphi[\sigma]$ iff $\mathcal{M} \not\models \varphi[\sigma]$;

(v) the structure satisfies $(\varphi \vee \psi)[\sigma]$ iff $\mathcal{M} \models \varphi[\sigma]$ or $\mathcal{M} \models \psi[\sigma]$;

(vi) the structure satisfies $(\varphi \wedge \psi)[\sigma]$ iff $\mathcal{M} \models \varphi[\sigma]$ and $\mathcal{M} \models \psi[\sigma]$;

(vii) the structure satisfies $(\forall x.\, \varphi)[\sigma]$ iff $\mathcal{M} \models \varphi[\sigma(m/x)]$ for all $m \in M$;

(viii) the structure satisfies $(\exists x.\, \varphi)[\sigma]$ iff $\mathcal{M} \models \varphi[\sigma(m/x)]$ for some $m \in M$.

If t is a closed term, then $t^{\mathcal{M}}[\sigma]$ is independent of σ, and we simply write $t^{\mathcal{M}}$. Similarly, if φ is a sentence, we often write $\mathcal{M} \models \varphi$ instead of $\mathcal{M} \models \varphi[\sigma]$, since the choice of σ does not matter.

2.2. Two-sorted first-order logic

Two-sorted first-order logic extends the single-sorted setting in a routine way. We will only record the differences and skip the analogous definitions. A systematic presentation can be found in [CN06, Section 4B; CN10, Section IV.2]. In principle one could work with arbitrary pairs of sorts, but in this thesis we instantiate the framework to the familiar number sort (ranging over $\mathbb{N}$) and string sort (ranging over finite binary strings). The goal of this section is therefore to emphasise what changes when we move from the definitions of Section 2.1 to this concrete two-sorted setting.

Definition 2.2.1 (Two-sorted first-order vocabularies)**.** A *two-sorted first-order vocabulary* (often abbreviated simply as a two-sorted language) $\mathcal{L}$ consists of collections of function and predicate symbols, much like an ordinary single-sorted vocabulary, but now the symbols may accept arguments of either of the two sorts. Moreover, the function symbols come in two varieties:

  (i) *number-valued* function symbols, whose outputs lie in the number sort; and
  (ii) *string-valued* function symbols, whose outputs lie in the string sort.

  For any pair $n, m \in \mathbb{N}$, the vocabulary contains:

  (i) a set of (n, m)-ary number-function symbols;
  (ii) a set of (n, m)-ary string-function symbols; and
  (iii) a set of (n, m)-ary predicate symbols.

A $(0, 0)$-ary function symbol is simply a constant symbol, which may be either a constant of the number sort or a constant of the string sort.

  We use $f, g, h, \ldots$ as metavariables for number-valued function symbols, $F, G, H, \ldots$ for string-function symbols, and $P, Q, R, \ldots$ for predicate symbols.

Definition 2.2.2 (The language $\mathcal{L}_A^2$)**.** As an example, consider the following two-sorted extension of the arithmetical language $\mathcal{L}_A$ (Definition 2.1.4):

$$\mathcal{L}_A^2 \; = \; [\, 0, \; 1, \; +, \; \cdot, \; |\cdot|\,;\; =_1, \; =_2, \; \leq, \; \in \,].$$

  Here the symbols $0, 1, +, \cdot, =_1, \leq$ are symbols of $\mathcal{L}_A$ (with $=_1$ corresponding to the usual equality of numbers). The symbol $\vec{X}$ is a number-valued function symbol giving the length of a string X. The binary predicate $\in$ relates a number and a string and is used to express membership: intuitively, $i \in X$ means that the i-th bit of the string X is 1. The symbol $=_2$ denotes equality between objects of the second sort.

  For convenience, when t is a number term, we abbreviate

$$X(t) \; \coloneqq \; t \in X.$$

Thus $X(i)$ plays the role of the i-th bit of the binary string X.

Remark 2.4. Note that $X(i)$ and $|X|$ intuitively should be related in some way. We don't set any implicit assumptions on the models we will consider. The necessary conditions will be fixed by the axiomatic system from Definition 9.2.1.

  In $\mathcal{L}_A^2$, the symbols $+$ and $\cdot$ each have arity $(2, 0)$; the length function $|\cdot|$ has arity $(0, 1)$; and the predicate $\in$ has arity $(1, 1)$.

Notation 2 (Bounded formulas). Let $\mathcal{L}$ be a two-sorted vocabulary. If x is a number variable and X a string variable that do not occur in the $\mathcal{L}$-number term t, we use the following abbreviations:

$$\exists x \leq t.\, \varphi \ \text{ stands for } \ \exists x.\,(x \leq t \ \wedge \ \varphi),$$

$$\forall x \leq t.\, \varphi \ \text{ stands for } \ \forall x.\,(x \leq t \ \rightarrow \ \varphi),$$

$$\exists X \leq t.\, \varphi \ \text{ stands for } \ \exists X.\,(\vec{X} \leq t \ \wedge \ \varphi),$$

$$\forall X \leq t.\, \varphi \ \text{ stands for } \ \forall X.\,(\vec{X} \leq t \ \rightarrow \ \varphi).$$

A quantifier appearing in one of these forms is called *bounded*, and a *bounded formula* is a formula in which every quantifier is bounded.

Notation. The expression $\exists \vec{x} \leq \vec{t}.\, \varphi$ abbreviates a block of bounded number quantifiers $\exists x_1 \leq t_1. \cdots \exists x_k \leq t_k.\, \varphi$ for some k, where no variable x_i occurs in any term t_j (even when $i < j$). The same convention applies to $\forall \vec{x} \leq \vec{t}$, $\exists \vec{X} \leq \vec{t}$, and $\forall \vec{X} \leq \vec{t}$.

Definition 2.2.3 (The $\Sigma_1^1(\mathcal{L})$, $\Sigma_i^B(\mathcal{L})$, and $\Pi_i^B(\mathcal{L})$ formulas). Let $\mathcal{L} \supseteq \mathcal{L}_A^2$ be a two-sorted vocabulary.
  (i) the class $\Sigma_0^B(\mathcal{L}) = \Pi_0^B(\mathcal{L})$ consists of all $\mathcal{L}$-formulas whose only quantifiers are *bounded number quantifiers* (string variables may occur free);
  (ii) for $i \geq 0$, the class $\Sigma_{i+1}^B(\mathcal{L})$ (resp. $\Pi_{i+1}^B(\mathcal{L})$) consists of formulas of the form

$$\exists \vec{X} \leq \vec{t}.\, \varphi(\vec{X}) \quad (\text{resp. } \forall \vec{X} \leq \vec{t}.\, \varphi(\vec{X})),$$

where:
   (i) $\vec{X}$ is a vector of string variables;
   (ii) $\vec{t}$ is a vector of $\mathcal{L}_A^2$-terms not involving variables from $\vec{X}$;
   (iii) φ is a $\Pi_i^B(\mathcal{L})$ formula (resp. a $\Sigma_i^B(\mathcal{L})$ formula).
  (iii) a $\Sigma_1^1(\mathcal{L})$ formula is a formula of the form $\exists \vec{X}.\, \varphi$, where $\vec{X}$ is a vector of zero or more string variables and φ is a $\Sigma_0^B(\mathcal{L})$ formula.
    We usually write Σ_i^B for $\Sigma_i^B(\mathcal{L}_A^2)$ and Π_i^B for $\Pi_i^B(\mathcal{L}_A^2)$.

Remark 2.5. The above definition might seem excessively limiting, as it doesn't allow us to have a string quantifier under a number quantifier. An example of a formula which is in none of Σ_i^B is $\forall x.\forall X.\exists Z.\, \varphi(x, Y, Z)$. The limitation of shape of formulas here is not a limitation on all the two-sorted formulas we can study. Indeed, we will reason about such formulas — we will just not say that they are in Σ_i^B. Details of one such reasoning are e.g. in [CN10, Lemma V.4.10; CN06, Lemma 5.35].

Remark 2.6. The above definition also doesn't allow using X in the term t used to "guard" a bounded quantifier $\exists X \leq t$. Note that if you used $t := X$, such a bounded quantifier would expand to $\exists X \leq X$, which is obviously equivalent to an unbounded quantifier $\exists X$.

Remark 2.7. The formalism described above is similar to the usual weak monadic second-order logic (WMSO) on words. We phrase it as a two-sorted first-order system to match the presentation in [CN06; CN10] and don't discuss the differences between the two formalisms here.

2.3. Decisional Turing machine complexity classes

In this section we introduce standard complexity classes such as L, P and NP. It is important to note that these classes only contain *decision* problems, only require i.e. computing a function $f : \{0,1\}^* \to \{0,1\}$.

Notation 3. We say that a machine *decides* a language $L \subseteq \{0,1\}^*$ iff it computes the function $f_L : \{0,1\}^* \to \{0,1\}$, where $f_L(x) = 1 \iff x \in L$.

Definition 2.3.1 (Time complexity [AB07, Definition 1.19; AB09, Definition 1.12])**.**
Let $T : \mathbb{N} \to \mathbb{N}$ be a function. We define $\mathsf{DTIME}(T(n))$ to be the class of all Boolean functions that are computable by some deterministic Turing machine running in at most $c \cdot T(n)$ steps on every input of length n, for some constant $c > 0$.

Definition 2.3.2 (Space complexity [AB07, Definition 4.1; AB09, Definition 4.1])**.**
Let $S : \mathbb{N} \to \mathbb{N}$ be a function and let $L \subseteq \{0,1\}^*$ be a language. We say that $L \in \mathsf{SPACE}(S(n))$ iff there exist a constant $c > 0$ and a deterministic Turing machine M deciding L such that, on every input of length n, the machine M visits at most $c \cdot S(n)$ distinct cells on its read–write work tapes (the input tape is read-only and does not count toward the space bound).

Definition 2.3.3 (Logarithmic space [AB09, Definition 4.5; AB07, Definition 4.5])**.**

$$\mathsf{L} = \mathsf{SPACE}(\log n).$$

Definition 2.3.4 (Polynomial time [AB09, Definition 1.13; AB07, Definition 1.20])**.**

$$\mathsf{P} = \bigcup_{c \geq 1} \mathsf{DTIME}(n^c).$$

Definition 2.3.5 (The class NP [AB09, Definition 2.1; AB07, Definition 2.1])**.**
A language $L \subseteq \{0,1\}^*$ belongs to NP if there exist

(i) a polynomial $p : \mathbb{N} \to \mathbb{N}$ (bounding the length of a certificate); and

(ii) a deterministic polynomial-time Turing machine M (called a *verifier* for L),

such that for every input string $x \in \{0,1\}^*$,

$$x \in L \iff \exists\, u \in \{0,1\}^{p(|x|)} \text{ with } M(x,u) = 1.$$

Whenever $x \in L$ and a string $u \in \{0,1\}^{p(|x|)}$ satisfies $M(x,u) = 1$, the string u is called a *certificate* (or *witness*) for x with respect to the language L and the verifier M.

2.4. Classes of binary circuits

We omit the standard definition of a Boolean circuit (a kind of finite directed acyclic graph of n inputs and m outputs). We also omit the definitions of circuit size, depth, and value; these are discussed in detail e.g. in [Vol99].

Definition 2.4.1 ($\mathsf{NC}^i_{/poly}$, $\mathsf{AC}^i_{/poly}$, $\mathsf{TC}^i_{/poly}$)**.** Fix $i \geq 0$. A language $L \subseteq \{0,1\}^*$ belongs to one of the following circuit classes if there exists a family of circuits $\{C_n\}_{n \in \mathbb{N}}$ such that $C_n(x) = 1$ iff $x \in L$ and:

(i) every circuit C_n has polynomially many gates w.r.t. n;

(ii) every circuit C_n has depth $\mathcal{O}((\log n)^i)$;

(iii) $L \in \mathsf{NC}^i_{/poly}$ (Nick's class) if each C_n uses only fan-in 2 $\wedge$-, fan-in 2 $\vee$-gates and fan-in 1 $\neg$-gates;

(iv) $L \in \mathsf{AC}^i_{/poly}$ (Alternating circuits) if each C_n uses unbounded fan-in $\wedge$-, unbounded fan-in $\vee$-gates and fan-in 1 $\neg$-gates;

(v) $L \in \mathsf{TC}^i_{/poly}$ (Threshold circuits) if each C_n uses unbounded fan-in $\wedge$-, unbounded fan-in $\vee$-gates, fan-in 1 $\neg$-gates and unbounded fan-in *majority* gates (i.e. a gate that outputs 1 iff at least half of its inputs are one).

Note that the condition $C_n(x) = 1$ iff $x \in L$ requires the circuits to have precisely one output node. We lift this requirement in Definition 2.4.4.

Definition 2.4.2 ($\mathsf{FNC}^i_{/poly}$, $\mathsf{FAC}^i_{/poly}$, $\mathsf{FTC}^i_{/poly}$). A function $F : \{0,1\}^* \to \{0,1\}^*$ belongs to $\mathsf{FNC}^i_{/poly}$(resp. $\mathsf{FAC}^i_{/poly}$, $\mathsf{FTC}^i_{/poly}$) iff there exists a family of circuits with multiple output nodes $\langle C_n \rangle_{n \in \mathbb{N}}$ such that whenever the input bits of C_n encode $X \in \{0,1\}^n$, the output bits encode $F(X)$; C_n satisfies the size and depth conditions (i), (ii) of Definition 2.4.1; and additionally C_n satisfies the class condition (iii) (resp. (iv), (v)) of 2.4.1.

Definition 2.4.3 (L-uniform circuit families). We say that a family of circuits $\langle C_n \rangle_{n \in \mathbb{N}}$ is L-uniform if there exists a Turing machine operating in space $\mathcal{O}(\log n)$, computing the function $1^n \to C_n$ for some representation of C_n.

Definition 2.4.4 (NC^i, AC^i, TC^i, FNC^i, FAC^i, FTC^i). A function $F : \{0,1\}^* \to \{0,1\}^*$ belongs to FNC^i(resp. FAC^i, FTC^i) iff there exists a L-uniform family of circuits satisfying the conditions from Definition 2.4.2.

Chapter 3

Models of computation, programming paradigms and complexity measures

One of the first decisions a programming language designer has to make is choosing the programming paradigm convenient for writing the programs of interest. In the practice of programming, imperative languages have no competition when a user needs to reason about the computational complexity of programs. The structure of imperative programs closely mirrors how the computation is executed on modern CPUs and GPUs. In turn, it is inherently unintuitive to reason about the complexity of programs written e.g. in Haskell or Prolog.

Yet if we want to understand what classes of functions can be characterized syntactically, we have to temporarily step away from the imperative mindset. As argued in [Gur12], the very notion of an "algorithm" is still evolving, so we shouldn't limit our considerations to a single paradigm. This chapter tests whether alternative models of computation could be better suited to form the basis of languages that capture popular complexity classes. This is studied in much more detail in [AB09, Section 1.6.3; AB07, Section 1.5.2], where also complexity of random computation, quantum computation and computation on *real* Turing machines is considered.

We can foreshadow that the answer is (perhaps surprisingly) positive: even though the complexity classes are defined on Turing machines, the characterizations studied in literature are rarely imperative.

3.1. Finite automata and transducers

Finite-state transducers compute string-to-string functions and have simple descriptions. In [Boj18], four characterizations are given for the class of *polyregular* functions — a class of string-to-string functions computed by a particular kind of transducer. The definitions described there readily constitute the basis of a programming language. Another programming language for transducers is studied in [Sch05]. A particular class of string-to-string functions defined using logic, *MSO transductions*, is characterized to be precisely the class of functions computable by two-way deterministic finite transducers (2DFT) in [EH99]. All of these results provide basis for new programming languages and are therefore more than relevant to our work.

Because transducer classes remain poorly understood, it is usually unclear whether an arbitrary problem belongs to the class recognized by a given flavor of transducer. Since writing

programs as transducers would rarely be useful for certifying an arbitrary program to be in the desired complexity class, for now we focused on other styles of characterization. A good overview of existing research on transducers can be found in [MP19].

The existing research on finite-state automata has not been directly useful for our work — perhaps because the expressiveness of this model is inherently limited to Boolean-valued functions. Yet, we consider some characterization of REG in Section 5.4.

3.2. Turing machines

This model of computation underpins the imperative programming style. There is a variety of flavors of Turing machines, and the details of a specific definition will most of the time not affect our considerations in this work. When not explicitly stating otherwise, when describing a computational process we will implicitly assume the realization of it on some kind of a Turing machine. As we will see in Chapter 6, the most popular complexity classes were specifically defined in terms of computation on Turing machines. The obvious approach to our task of obtaining certificates of complexity bounds would be to start from the bare definition of a Turing machine, and meticulously formally proving properties of more and more complex Turing machines. Yet, this approach fails miserably, as we will have a chance to explore in Chapter 4.

Surprisingly few elegant, high-level imperative languages are known to characterize well-studied complexity classes beyond the trivial ones. This scarcity is precisely why much of the work surveyed in later chapters relies on non-imperative paradigms.

3.2.1. Random-access Turing machines

Random-access Turing machines are Turing machines with a special "pointer" tape, of length logarithmic to the size of input, and a special state such that when the binary number on the pointer tape is n, n-th digit of the input is written to the work tape.

Reasoning about computation in complexity classes such as L or P is the same for traditional Turing machines and random-access Turing machines. The choice of model starts to matter for notions of *fine-grained* complexity classes such as $\mathsf{DTIME}(n)$ Definition 2.3.1 or $\mathsf{DTIME}(n^2)$. Due to insufficient existing research on implicit characterizations of fine-grained complexity classes, we will not consider them besides a brief discussion in Subsection 6.3.2. The notion of $\mathsf{DLOGTIME} = \mathsf{DTIME}(\log n)$-uniformity explored in **??** is also only defined for random-access Turing machines.

3.3. Circuits

Circuits as a model of computation are the theoretical foundation of parallel programming. The theoretical implications turn out to not be widely applicable in practice, however. In this work, we will mostly use circuits to reason about very weak complexity.

A very rich overview of different characterizations of circuit complexity classes is in [ADK25]. A particular logical characterization of AC^0 (Definition 2.4.4) will be important for us in Chapter 9.

3.4. Discrete differential equations

An original point of view on computation is to describe functions as solutions to discrete differential equations. For example, in [BD19], FP (Definition 6.2.2) and FNP (Definition 6.2.3) are characterized. Characterizations of various circuit complexity classes from FAC^0 to FAC^1 (Definition 2.4.4) are described in [ADK25].

3.5. Logic programming and descriptive complexity

Logic provides very deep complexity-theoretic connections, primarily through descriptive complexity theory, which we explore in more detail in Chapter 5.

3.6. Untyped recursion

Another classical paradigm is that of general recursive functions, or equivalently the untyped lambda calculus. Because these systems are Turing complete, the interesting question is how to constrain recursion so that the resulting language captures a specific class. We treat these ideas in Chapter 7.

TODO: function algebras?

3.7. Typed lambda calculus

Typed lambda calculus underpins functional programming. In this section we focus on typed variants, unlike in Section 3.6.

Lambda calculus does not line up cleanly with traditional Turing machine complexity measures. For instance, for some representation of strings $\{0,1\}^*$, [HK96, Theorem 3.4] identifies the functions $\{0,1\}^* \to \{0,1\}$ definable in simply typed lambda calculus (STLC), with regular languages. But with a different encoding of inputs, [HKM96] relates STLC to the whole ELEMENTARY class. Moreover, [Zak07] states that with a different "standard" encoding, STLC instead characterizes extended polynomials, and further shows that if we slightly modify the encoding, the class is yet different. For more discussion, see also [Maz18]. Consequently, it is not obvious how to reason about complexity theory in the language of lambda calculus.

Nevertheless, typed lambda calculi have been utilized very successfully to syntactically characterize complexity classes. Recall that one of the reasons linear logic is studied is the potential to reason about resource creation and utilization. Concepts from linear logic have been implemented in the theory of type systems to transfer the resource interpretation. This will be discussed further in Chapter 8.

3.8. Set theory as inspiration for model of computation

An interesting connection appears when we think of traditional notions of "complexity" of sets in set theory from the point of view of computational complexity. If we treat taking complements, intersections, countable unions as operations in a programming language, perhaps we could design a programming language for constructing sets. By itself such a language would probably not be the most interesting one. However, thinking about mathematical reasoning in terms of a computational process is a very powerful technique. It has been particulary deeply explored for connecting logic and lambda calculus under the name of Curry-Howard or proofs-as-programs correspondence.[1] For the computational content of set theory in particular, an interesting brief discussion is presented in [Tao10].

For our purposes, an interesting connection appears when we look at descriptive set theory. The most basic classes of sets distinguished there are open and closed sets. Slightly higher, an F_σ set is a countable union of closed sets; a G_δ set is a countable intersection of open sets. A few levels higher up the *Boldface hierarchy*, which in a way quantify the complexity of sets, Borel sets are considered:

Definition (Borel sets)**.** Let X be a topological space. The class of Borel sets of X, $\mathcal{B}(X)$ is the smallest class of sets containing every open set of X and closed under (i) and (ii):

(i) if A is Borel, then its complement $X \setminus A$ is Borel;

(ii) if A_n is Borel for each $n \in \mathbb{N}$, then the countable union $\bigcup_{n \in \mathbb{N}} A_n$ is Borel.

As this is a standard least fixed point definition, it is strikingly similar to definitions of classes of recursive functions that we will consider later, e.g. Definition 7.1.1. We can also take a computational point of view on theorems about the determinacy of Gale-Stewart games (which we shall not introduce here). It is widely known that Gale-Stewart games are determined when the underlying set is open or closed. Allowing the underlying set to be more and more complicated, we quickly reach the limits of provability in Zermelo-Fraenkel set theory: determinacy for Borel sets is a difficult theorem, and determinacy for analytic and projective sets is independent of ZF, yet provable assuming as axiom the existence of an appropriately large cardinal.

A good question to ask is if by carefully curating the axioms, we could obtain a mathematical theory such that theorems corresponding to computation in our desired complexity class are provable, and the theorems that wouldn't be "implementable" are not.

We will circle back to this intuition while considering the `PIGEON` computational problem in 6.2.2 and the unprovability of the related pigeonhole principle in the weak theories studied in 9.9. Especially the last fact about unprovability is interesting for us in this section, as it resembles e.g. independence of continuum hypothesis from ZFC, but in a strictly computational setting of not being able to perform enough computation in a low complexity class.

> TODO: (optional) The forcing technique used originally by Cohen to prove the mentioned independence, turns out to also be useful for the theory of computation, as discussed in Subsection 6.3.1.

While this line of thought is very far from "practical programming languages", these considerations inspired[2] the approach that we study in Chapter 9.

---

[1] A good introduction to the immensely deep topic of proofs-as-programs is [SU06].

[2] This line of study grew out of presenting the topic at the JAiO master's seminar at the University of Warsaw — slides are available online at [Bal25].

Chapter 4

Formalized semantics

One approach to certify complexity class of a program would be to simply accept an implementation of the algorithm in C++ accompanied by a proof that the number of "steps" of the algorithm when executed on a standard computer is bounded by a polynomial p. For at least three reasons, such a language would not be satisfactory for us. We explore them in this chapter.

TODO: careful with three

Informality of language semantics It is not actually defined how much time will an arbitrary C++ program take to execute. A famous shortcoming of C++'s standard is the notion of *undefined behaviour*. As it turns out, close to none programming languages have well-defined semantics. For the vast majority of languages, the only reliable semantics is the source code of the most popular compiler. There are, however, programming languages with formally defined semantics. Most notably, a large fragment of the C programming language has been formalized in the CompCert project [Ler09]. CakeML [Kum+14] has formal semantics and a proven-correct compiler. A fragment of OCaml has been formalized in [Sea+25].

Representation of a proof in computer-verifiable form Even if someone gives us a program with a well-defined semantics and the proof required, it is not obvious how to check that the proof is correct. A system in which it is convenient to write mathematical proofs, in a way that they can be checked by the computer, and the code of the checker is simple enough to be widely believable, is called a *proof assistant*. The successful projects such as Mizar, Rocq, Lean and Isabelle/HOL took decades of work to be developed. Only in the last decade, the most popular proof assistants got to the stage where research-level reasoning can be transferred to them *with reasonable overhead*.

In a separate but related field, a very interesting class of programming languages rely on automated theorem proving to verify user-provided pre- and post-conditions of functions. Dafny [Lei10] (see: Listing 4.1) does it by utilizing *SMT-solvers*, not allowing the user to conduct a mathematical proof manually. Why3 [Bob+11] provides a comprehensive framework, combining the power of SMT-solvers for automation where possible, and allowing the user to conduct proofs manually where they fail.

Infeasibility of formal proofs about Turing machines Even while operating in the right programming language, and assuming trust in the modern proof assistants, we will still not be able to proceed this way. The field of computational complexity and computability theory has a reputation of being particularly hand-wavy. This naturally generates difficulties when attempting to formalize the proofs from these areas. Voluminous discussion on that problem

```dafny
function fib(n: nat): nat
{
  if n == 0 then 0
  else if n == 1 then 1
  else fib(n - 1) + fib(n - 2)
}
method ComputeFib(n: nat) returns (b: nat)
  ensures b == fib(n)
{
  var i := 1;
  var a := 0;
  b := 1;
  while i < n
    invariant 0 < i <= n
    invariant a == fib(i - 1)
    invariant b == fib(i)
  {
    a, b := b, a + b;
    i := i + 1;
  }
}
```

Listing 4.1: Dafny example

emerged with the rise of popularity of proof assistants. The below quote is due to Yannick Forster, author of the most successful works on formalizing results about computability:

> making these arguments [i.e. about Turing machines] formal is several orders of magnitude more involved than formalising other areas of mathematics, due to the amount of invisible mathematics (a term coined by Andrej Bauer) involved. [For25]

Formally defining Turing machines has mostly been considered in [AR12] and, building on top of it, [FKW20][1]. Besides that, few examples of concrete Turing machine definitions have been described in the literature, as working with them directly is greatly time-consuming and not incentivized by academia. Interesting concrete examples were provided in[Kud96], [Rog96] and in the work we discuss in Section 4.1. Alternative approaches are studied as Synthetic Computability [Bau06].

In Section 7.3.1, we briefly discuss one formalization of *a characterization* of polynomial-time functions.

4.1. Ciaffaglione's formalization of undecidability of HALT

In [Cia16], an elegant formalization of the undecidability of the halting problem is presented. The proof goes by defining simple Turing machines needed and directly proving results about their semantics.[2] As part of our work, we refined the code to a newer Coq version, fixing the

---

[1]Part of an ongoing formalization project: `https://github.com/uds-psl/coq-library-undecidability`.
[2]As of November 2025, the code referred to in the paper is not accessible anymore.

proofs where necessary. Full code (also presented on the JAiO seminar) is available online
with installation and verification instructions.[3]

TODO: ensure code didn't move

---

[3]Code and slides are available at: `https://github.com/ruplet/prezentacja-seminarium-1`.

Chapter 5

Descriptive Complexity

In the rest of this book, we usually measure complexity of algorithms: how much time and memory will a given algorithm need to run to solve a problem of size n? In this chapter we will instead focus on the complexity of defining the problem itself.

The central problem of Descriptive Complexity is to characterize a complexity class by the *power of logic required to define its problems*. This is in contrast to *Implicit Complexity Theory*, discussed later in Section 7.3 and Section 8.1, which seeks the weakest system sufficient to *implement algorithms* of the class; and in contrast to *Bounded Arithmetic*, which studies the weakest *theory* required to define a function *and prove its correctness*. This perspective has led to elegant logical characterizations of many traditional complexity classes, as we will discuss in Section 5.1.

Before going into details, we will discuss an example to recall what it means for a structure to model a formula (Definition 2.1.7):

Remark 5.1. In Descriptive Complexity, there is no notion of a proof. The problem of deciding if a given sentence is provable is independent of this chapter's considerations, and will be studied by us only in Chapter 9.

Example 5.2. Consider a vocabulary $\mathcal{L}$ consisting of unary relations $\mathsf{Zero}(x), \mathsf{One}(x)$ and binary relations $=, \leqslant$. If we think of positions in a binary string as elements of the universe, we can define e.g. for a string 01011 an $\mathcal{L}$-structure $\mathcal{M}$ with:

(i) universe $\{1, 2, 3, 4, 5\}$;

(ii) $\mathsf{Zero} := \{1, 3\}; \mathsf{One} := \{2, 4, 5\}; = := \{\langle 1, 1 \rangle, \langle 2, 2 \rangle, \ldots, \langle 5, 5 \rangle\}; \leqslant := \{\ldots\}$.

Now, we can reason about the original binary string using logic. The formula $\exists x. \mathsf{One}(x)$ is true in $\mathcal{M}$. However, the formula $\forall x. \mathsf{Zero}(x)$ is not. It is insightful to analyze in general what kind of logical formulas are true in $\mathcal{M}$.

> TODO: keep 0- or 1- indexing consistent with rest

> TODO: is this insightful or another question?

We can use the language of logic to describe computational queries about the underlying structure. This is the intuition behind the languages designed for querying databases. To formally connect logic and computation, consider the following

Definition 5.0.1 (Language of binary strings)**.** Define the vocabulary $\tau_{\leqslant}$ to contain only the unary relations $\mathsf{Zero}(x), \mathsf{One}(x)$ and the binary relations $=, \leqslant$. Intuitively, $x \leqslant y$ means that memory cell x comes before memory cell y. As we can represent most of the real-world structures of interest as some binary string on a computer, this simple vocabulary already allows us to ask interesting questions.

Definition 5.0.2 (First order logic on ordered structures)**.** $\mathsf{FO}_{\leqslant}$ is the class of sentences of first-order logic in the language $\tau_{\leqslant}$.

Remark 5.3 (Bibliography). In literature, first-order logic over $\tau_\leqslant$ is called $\mathsf{FO}(\mathsf{wo}\,\mathsf{BIT})$; see [Imm99, Ordering Proviso 1.14].

Definition 5.0.3 (Language generated by a formula). For a binary string $w \in \{0,1\}^n$ define its corresponding structure $\mathcal{M}_w$ over vocabulary $\tau_\leqslant$ to have the universe $M := \{0, 1, \ldots, n-1\}$ and the standard semantics of relations from $\tau_\leqslant$.

We define the language L_ϕ to be *generated* by a first-order sentence $\phi \in \mathsf{FO}_\leqslant$ iff

$$\mathcal{M}_w \models \phi \iff w \in L_\phi.$$

Let's also define a logic, where elements of the universe are treated not as memory cells, but as natural numbers.

Definition 5.0.4 (First order logic with arithmetical predicates). $\mathsf{FO}_{\mathsf{BIT}}$ is the class of sentences of first-order logic with the following vocabulary:

  (i) constants $0, 1, \mathsf{max}$;
  (ii) binary relations $\mathsf{BIT}, \mathsf{SUCC}, \leqslant, =$.

From now on, whenever talking about the semantics of any such sentence, we will require the elements of the universe to be interpreted as actual natural numbers $0, 1, \ldots$; the equality and order to be interpreted as the actual equality and order on the elements of the model; $0, 1, \mathsf{max}$ to be interpreted as the minimum, second and maximum elements under $\leqslant$; $\mathsf{BIT}(x,y)$ to hold iff bit y in the binary representation of x is 1.

Remark 5.4. Note that it is a completely different definition to Definition 5.0.2. However, as hinted in Remark 5.3, the expressive power of $\mathsf{FO}_{\mathsf{BIT}}$ without BIT for our purposes is the same as $\mathsf{FO}_\leqslant$ and we will not discuss these subtleties here.

<div style="border: 1px solid #ccc; padding: 4px; background: #fde8d0;">TODO: tutaj blefuję, nie mam cytowania na to ze to jest to samo</div>

5.1. Results

The field of descriptive complexity has a long history and we should not introduce all the results in this work, as our primary purpose is to examine if, in the first place, they are relevant to our problem. We will only describe the logical characterizations of the classes that are the most interesting for us: P and uniform AC^0 which will underpin our Chapter 9. A concise overview of the classical results is presented in [Imm99, Section 15.1], where characterizations of $\mathsf{SPACE}(n^k)$, L, NL, P, NP and PSPACE are described.

<div style="border: 1px solid #ccc; padding: 4px; background: #fde8d0;">TODO: AC0 here is fo-uniform AC0 which is good. In my preliminaries, i set AC0 logspace-uniform, which is almost as strong as logspace. fix that in preliminaries</div>

Theorem 5.5 ([Imm99, Corollary 5.32] $\mathsf{FO}_{\mathsf{BIT}} = \mathsf{FO}\text{-uniform } \mathsf{AC}^0$). *Languages $L \subseteq \{0,1\}^*$ decidable by uniform AC^0 circuits are precisely the languages generated by FO formulas, in the sense of Definition 5.0.3.*

Remark 5.6 (Bibliography). The theorem is originally stated in terms of $\mathsf{FO}_{\mathsf{BIT}}[\mathsf{t(n)}]$ defined in [Imm99, Definition 4.24]. We limit to $t(n) = \mathcal{O}(1)$, i.e. in our case $\mathsf{FO}_{\mathsf{BIT}}[\mathsf{t(n)}] = \mathsf{FO}_{\mathsf{BIT}}$. The uniformity condition is also different to the one we fixed in Definition 2.4.3: their circuits are so-called FO-uniform, which is a much stronger condition. For details about the different notions of uniformity, refer to Appendix A.

5.2. The Quest for a Logic Capturing PTIME

Generalizing our considerations, we may wonder what would happen if we removed the $\leqslant$ relation from the vocabulary. The difference turns out to be an object of an interesting branch of research. One intensely studied logic is $\mathsf{FO}[\mathsf{LFP}]$, defined inductively.

Definition 5.2.1 ([Imm99, Definition 4.5] Least fixed-point logic)**.** FO[LFP] is a class of logical sentences such that:

(i) if $\varphi \in$ FO then $\varphi \in$ FO[LFP];

(ii) if $\varphi_R(x_1, \ldots, x_k) \in$ FO[LFP], where R is a k-ary relation and R only occurs positively in $\varphi_R(x_1, \ldots, x_k)$ (i.e. every occurrence of R is preceded by an even number of negations), then LFP$_{R(x_1,\ldots,x_k)}\varphi$ may be used as a new k-ary relation symbol denoting the least fixed-point of φ.

By a result independently proved by Immerman [Imm86] and Vardi [Var82], the class of languages generated by sentences from FO[LFP] is precisely P [Imm99, Theorem 4.10].

Without an explicit ordering (i.e. on *unordered structures*), FO[LFP] does **not** capture P [CFI92]. The existence of a logic that characterizes P on unordered structures remains a major open problem in computer science as of 2025. The complexity class which would be captured by such a logic is invP, the class of permutation-invariant polynomial-time problems. In other words, invP is the class of polynomial-time boolean functions that take a graph on input and have the property that the result is the same for all orderings of the graph nodes. A good overview of this problem is [Daw12] and [Imm99, Chapter 12, The Role of Ordering].

An important result that treats unordered structures is Fagin's theorem [Fag74] that states that the class of languages generated by sentences of existential second-order logic is precisely NP. Intuitively, ordering of the domain is not necessary to assume for Fagin's theorem because in NP, we can *guess* it.

5.3. Defining functional problems in logic

5.3.1. First-order queries (FO-reductions)

Despite logic being only able to naturally define decisive problems, some approaches have been used to reason logically about functions. Most importantly, to study completeness of problems in low complexity classes such as L, FO-reductions are used. They have a rather complicated definition, which we don't display here and for the details refer to [Imm86, Definition 1.26]. In the same work, even weaker notions of reducibility are studied: first-order projections (fops) and quantifier-free projections (qfps) are defined in [Imm99, Definition 11.7]. An interesting property of the complexity classes we are studying in this work is that their complete problems are already complete under surprisingly weak reductions. In [Imm99, Proposition 11.10] it is proved how SAT is NP-complete via fops and even via qfps.

5.3.2. Bit-graph definitions

An easy way to define a general function from Boolean functions is to use the Boolean functions to decide if "i-th bit of the output $f(x)$ is 0 or 1". For the definition to be precise, two more technicalities are needed (the output's length itself must be polynomial and easy to compute), which we will discuss while defining L-reductions in Definition 6.1.1. For now, we will just say that this style of definitions is very important for some of the results we will study in Chapter 9, as the "semantic" definition of definability of functions ([CN10, Definition V.4.12; CN06, Definition 5.37]) is precisely that.

5.3.3. FO and MSO transductions

Some works use the notion of FO-transductions, e.g. in [NMS21, Section 2], where they are also defined. They are, however, completely different anything we study in this work and

are defined as compositions of *copying*, *coloring* and *simple interpretations*, which we will not discuss. A similar, and much more important notion is of MSO transductions defined e.g. in [Cou94, Section 2]

5.4. Descriptive Complexity and programming languages

From the point of view of programming languages, the meaning of e.g. the result discussed in Section 5.2 is as follows: given a logical formula $\varphi \in$ FO[LFP], we *know that there exists some* Turing machine of complexity P that will check for us if $w \models \phi$ for any input w. There is a huge leap of faith, however — just that the machine *exists* and runs fast, we can't conclude that we will ever be able to actually use it. We have to inspect the proof of such a theorem and tell if we can compute the description itself of such a machine, and how fast we can do it. An illustrative example is infeasibility of using the below theorem to design a programming language for finite automata:

Theorem 5.7 ([Imm99, Theorem 1.36] Büchi-Elgot theorem). *For the alphabet $\Sigma := \{0, 1\}$, the set of boolean queries expressible in second-order monadic logic (which we skip the definition of) over the vocabulary $\tau_{\leqslant}$ consists exactly of the regular languages. In other words,* MSO = REG.

Remark 5.8 (Bibliography). The theorem was originally proved by Büchi in [Büc90] and by Elgot in [Elg61]. The statement from [Imm99, Theorem 1.36] is slightly more accessible. It uses slightly different wording to ours. For the details of the definitions, refer to [Imm99, Proviso 1.14] and [Imm99, Proviso 1.15].

However, not always is this problem that difficult. Results from descriptive complexity has been crucial for the field of database theory, which we don't discuss here. The approach of using logic for programming has also coined the paradigm of logic programming.

5.4.1. Logic programming

The most famous example of a logic programming language is Prolog. Prolog doesn't have a reputation of a language well-suited for computational complexity analysis. However, it relates very well with the notions of deciding computational problems we discussed in this chapter. Please see Listing 5.1 for a demonstration how can we transfer our considerations from Example 5.2 to practical computation.

5.4.2. Datalog

Datalog is a programming language that uses the paradigm of logic programming, similarly to Prolog. Unlike Prolog, however, it *captures* the complexity class of P [Dan+01, Theorem 4.4] (on ordered, finite structures). That is, it is a language such that any query definable in it will be checkable in P for a given model. We have not investigated the subtleties of Datalog's implementation and whether the complexity doesn't blow up somewhere.

5.4.3. Logic as the type system

Could the results from descriptive complexity be used to design a specification mechanism (such as a type system[1]) for a programming language? While such a language would be very interesting and possibly have highly desired properties, our "typechecker" in the naive case

---

[1]Language does not have to be typed to foster partial proofs of correctness; see e.g. [LP99], [Pau00]

```prolog
% Run at: https://swish.swi-prolog.org/
succ(bit1, bit2).
succ(bit2, bit3).
succ(bit3, bit4).
succ(bit4, bit5).

zero(bit1).
zero(bit3).
one(bit2).
one(bit4).
one(bit5).

exists_1  :- one(X).                           % true
exists_00 :- succ(X, Y), zero(X), zero(Y).     % false
no_00     :- \+ ( succ(X, Y), zero(X), zero(Y) ). % true
no_11     :- \+ ( succ(X, Y), one(X), one(Y) ).   % false
```

TODO: poprawic ta jedna linijke %true rozjechana

Listing 5.1: Prolog example

would essentially be a proof checker for first order logic. This would be a very difficult system to properly properly design. We were not successful in curating such a small set of basic programming instructions that, given a pre- and a post-condition in FO and an operation from that set, checking for the validity of such a triple would be an easy problem.

Nonetheless, this line of inquiry led to a broader investigation of type systems that enforce resource bounds — particularly those inspired by linear logic and implicit computational complexity. The results of this exploration are presented in Chapter 8.

Chapter 6

Reductions

Plenty of theorems of the form "problem P is *complete* for class C under reductions in class R" have been described in the literature. In this chapter we analyze when such reductions help capture complexity classes syntactically and when they fall short.

Definition 6.0.1 (Circuit Value Problem (CVP))**.** Given a representation of a Boolean circuit on input, compute its output.

> TODO: more formality

Theorem 6.1 ([Lad75])**.** *The circuit value problem is complete for* P *under* L *reductions.*

Remark 6.2. We define L-reductions formally in Definition 6.1.1.

    Knowing Theorem 6.1, we could design a language for P in such a way: first, design a language for L which is (perhaps) simpler; then, as the compiler provider, include a standard library function P, which the L-functions from the language could call like an oracle to get solution for P. In this chapter we try to understand if such a language would truly have the full power of polynomial-time functions. The core of this chapter is Theorem 6.9, stating that this holds at least for the class of logspace-computable functions and a particular notion of weak reductions. But the heart of this intuition is not developed until **??**, where we get such characterizations for most of the complexity classes that we consider in this thesis.

6.1. Decisional and functional complexity classes

We can't answer the question of expressive power of CVP + L-reductions by comparing it with any decisional class. Focusing solely on the complexity of Boolean functions as we did for now e.g. in Section 2.3 is not sufficient to reason about general functions with output. In this chapter we will introduce the still standard, but less talked about, *functional* complexity classes, that study general functions $f : \{0,1\}^* \to \{0,1\}^*$. In complexity theory these are usually implicitly used to define *reductions*. The results about these classes transfer much better to our interests than results about the decisional complexity classes. As we discuss in Chapter 7, programming-language-like characterizations of decisional complexity classes are far more abundant and predate the characterizations of functional complexity classes. It is the latter, however, that is viable for our purposes of designing a programming language.

    A thorough overview of complexity classes of functions is described in [Sel94]. A more thorough discussion of decision vs search is in [BG92].

Exponential-length output   The difference between P and FP is obvious when we look at the below example:

Example 6.3. Running time of any Turing machine computing the function $x \to 2^x$ for input and output in binary, is exponential. At the same time, given an input x, y, checking if $y = 2^x$ is easily in polynomial time.

This problem is, however, usually artificially mitigated by requiring the length $|f(x)|$ of the function's output to be polynomially bounded as in Definition 6.1.1.

Self-reducibility   A typical, and ubiquitous in the literature way of defining function problems is to require an external proof that $|f(x)|$ is polynomially bounded, then repeatedly decide if i-th bit of $f(x)$ is 0 or 1. For example, let's look at the below

Definition 6.1.1 ([AB09, Definition 4.16; AB07, Definition 4.14] L-reductions)**.** A function $f : \{0,1\}^* \to \{0,1\}^*$ is said to be *logspace computable* if

(i) the function f is *polynomially bounded*, meaning that there exists a constant $c > 0$ such that
$$|f(x)| \le |x|^c \qquad \text{for all } x \in \{0,1\}^*;$$

(ii) the following two languages lie in L:
$$L_f = \{\, \langle x, i \rangle \mid f(x)_i = 1 \,\}, \qquad L'_f = \{\, \langle x, i \rangle \mid i \le |f(x)| \,\}.$$

In other words, a deterministic $\mathcal{O}(\log|x|)$-space machine can, given (x, i), determine whether i is within the length of $f(x)$ and, if so, whether the i-th bit of $f(x)$ is 1.

A language B is *logspace reducible* to a language C, if there exists a function $f : \{0,1\}^* \to \{0,1\}^*$ that is logspace computable and such that $x \in B$ iff $f(x) \in C$ for every $x \in \{0,1\}^*$.

Famously, we can also apply this trick to solve the problem FSAT, the corresponding functional problem to the SAT of finding a specific satisfying assignment with polynomially many calls to the decision procedure SAT. In general, many functional problems are solvable in polynomial time with polynomially many calls to their corresponding decisional problems. We say that problems with that property are *self-reducible*.

However, some search problems are unlikely to be self-reducible. A good example is the problem of integer factorization, which is still, as of November 2025, conjectured to not be in P even despite the breakthrough result from 2002 in which PRIME (decide if n is prime) was proved to be in P [AKS04]. A particularly important class of such problems is considered in 6.2.2. But first, let's define the classes FP and FNP.

6.2. Functional complexity classes

Definition 6.2.1 ([AB09, Section 17.2; AB07, Section 9.1] The class FP (Version 1))**.** FP consists of all functions
$$f : \{0,1\}^* \to \{0,1\}^*$$
that are computable by a deterministic polynomial-time Turing machine. In contrast to decision problems (which output a single bit), functions in FP may produce outputs of arbitrary polynomial length.

Remark 6.4. This definition is used e.g. in [Coo85] besides the work cited above. It is rather equivalent to Definition 6.2.2 below, also ubiquitous in the literature.

Definition 6.2.2 ([Ric07, Section 28.10] FP (Version 2))**.** A binary relation $P(x, y)$ is in FP iff there exists a polynomial-time Turing machine that, given an arbitrary input x:

  (i) outputs some y such that $P(x, y)$ if any exists;

  (ii) signals that no such y exists otherwise.

Remark 6.5. This version can make it more obvious to compare FP with FNP (defined later) — but only assuming that FNP is defined using nondeterministic Turing machines, which is not true in our case; we will use the *verifier*-style definition.

Remark 6.6 (P vs FP)*.* These two classes are often identified due to similar properties. The notion of completeness for both of them, despite being differently defined, practically is practically the same to the robustness of P computations under being repeated for every bit of the output. Indeed, even in Stephen Cook's 1982 ACM Turing Award lecture [Coo83, Section 6], it is not clearly distinguished between P-completeness and FP-completeness: the 3 proofs cited in this lecture as proofs of FP-completeness of some functions $f(x)$ only themselves prove the P-completeness of problems of the form "decide if i-th bit of the result $f(x)$ is zero".

The two classes, however, are not the same. In [Kre88, Theorem 4.1], it is proved that if $\mathsf{FP}^{\mathsf{SAT}}[\mathcal{O}(\log n)] = \mathsf{FP}^{\mathsf{SAT}}[n^{\mathcal{O}(1)}]$ then also $\mathsf{P} = \mathsf{NP}$. In turn, as noted in [Har, discussion after Theorem 8], the corresponding result for P^{NP} versus $\mathsf{P}^{\mathsf{NP}}[\mathcal{O}(\log n)]$ is not known, and indeed fails relative to some oracles.

For a good discussion specifically on FP-completeness, which is relatively hard to find, there is an argument that finding the lexicographically first maximal clique in an undirected graph is NC^i-complete for FP in [Coo85, Proposition 6.1].

6.2.1. FNP

The definition of FNP is tricky to get right. A very good discussion of the awkwardness of the definitions is present in [httb]. For extensive discussion on the different definitions, see [httc], [htta]. In Papadimitriou's book, it's defined in yet another way, as a class function problems for NP, not in terms of a specific computational model.

Definition 6.2.3 (The class FNP)**.** A binary relation $P(x, y)$ is in FNP if there exist:

  (i) a polynomial $p : \mathbb{N} \to \mathbb{N}$ such that if for a given x exists a solution y such that $P(x, y)$, then there also exists a "short" solution y' such that $P(x, y')$ and $|y'| \leqslant p(|x|)$;

  (ii) a deterministic polynomial-time Turing machine M (a *verifier*), such that for every input pair (x, y),

$$P(x, y) \iff M(x, y) = 1.$$

Remark 6.7. This definition is in style of [Ric07, 28.10 and Theorem 28.9], where also the other, nondeterministic Turing machines-based definition is listed.

The other definition might come off as more intuitive: that a relation P is in FNP iff there is a nondeterministic polynomial-time algorithm that, given an arbitrary input x, can find some y such that $P(x, y)$ or signal that it doesn't exist [BD19]. However, as such nondeterministic Turing machines don't seem to be physically realisable, we don't want to introduce that computational model in this work.

6.2.2. NP vs FNP and the total search problems

Definition 6.2.4 (TFNP)**.** A binary relation $P(x, y)$ is in TFNP (total FNP) iff it is in FNP and for every x exists at least one y such that $P(x, y)$.

An interesting example of a problem in TFNP is PIGEON defined below, for which we mathematically know that the answer exists, but finding it is not trivial.

Definition 6.2.5 (PIGEON). Given a binary string encoding a Boolean circuit $C : \{0,1\}^n \to \{0,1\}^n$, return either an input x such that $C(x) = 0^n$, or two distinct inputs $x \neq y$ such that $C(x) = C(y)$.

Remark 6.8. This problem will be of our interest in Theorem 9.9, where we will discuss mathematical theories so weak that the pigeonhole principle is not their theorem. The intuition behind it is that the computational content of these theories is not strong enough to perform an exhaustive linear search of the whole domain.

The class PPP, a subclass of TFNP problems for which the solution is guaranteed to exist by the pigeonhole principle, is conjectured to not be equal to FP. If PPP = FP, then one-way permutations do not exist [Pap94, Proposition 3], which would have tremendous implications for cryptography.

The class TFNP is discussed in yet more detail in [Mor05, Section 1.1].

6.2.3. Language for FL

> TODO: I really don't want to introduce circuit reductions... skipping it.

Theorem 6.9 ([Coo85, Proposition 4.1]). *We obtain precisely the class* FL *from the closure of* L *under* NC^1 *circuit reductions, symbolically:* $\mathsf{FL} = \mathsf{L}^*$

Remark 6.10. Originally, this theorem is proved with NC^1-reducibility meaning reducibility by U_{E^*}-uniform NC^1 circuits. We don't introduce these notions in this work (except for a brief discussion of U_{E^*}-uniformity in Section A.2)

An overview of problems complete for L is present in [CM87].

6.2.4. Language for FP

We can derive an analogous result to Theorem 6.9 for the class FP. We are discussing it here, postponing the considerations to **??**.

The notion of P-completeness is defined formally e.g. in [AB07, Definition 6.25; AB09, Definition 6.28]. A very detailed description of one problem complete for P under L-reductions is in [Koz06].

6.2.5. Not-a-Language for FNP

There is a relatively agreed-upon notion of reductions between FNP problems:

Definition 6.2.6 ([Göb11; Gol21] Polynomial-time reductions for FNP). Let `HardProblem`, `NewProblem` be search problems in FNP. We say that `HardProblem` (many-one) reduces to `NewProblem` if there exist f, g in FP such that:

$$\texttt{NewProblem}(f(x), y) \implies \texttt{HardProblem}(x, g(y))$$

For a given input x of `HardProblem`, we can run `NewProblem`$(f(x))$ to obtain some result y, such that $g(y)$ is the result of `HardProblem`(x).

There is also plenty of NP-complete problems described in the literature. However, it is very unclear if we will get FNP this way. The class $\mathsf{FP}^{\mathsf{NP}}$ is well-studied and nothing suggests it to be equal to FNP.

6.2.6. Semantic and syntactic complexity classes

Some of the complexity classes remain notoriously difficult to be characterized implicitly, by e.g. showing a complete problem and reductions for it. However, as it turns out, not all complexity classes studied have known complete problems. The classes for which a complete problem exists are called "syntactic" complexity classes, as opposed to "semantic", e.g. in [GP18], the authors define a new complexity class $\mathsf{PTFNP} \subseteq \mathsf{TFNP}$, for which they prove the existence of a complete problem, and then call this class "syntactic".

An interesting discussion of this problem, centered around the class $\mathsf{inv} - \mathsf{P}$ we discussed in Section 5.2, is present in [Daw12]. For the class BPP, some discussion is in [httd]. Despite that, a "less implicit" characterization of BPP was studied in [LT12].

Interestingly, PP has been characterized implicitly by Ugo Dal Lago: [DKO21].

Remark 6.11 (Bibliography). For probably the first published recognition of the widespread inconsistency of decisional vs functional complexity classes in the literature, with examples of inconsistent places see [Lee91, Page 131] (and our bibliographical Remark A.1).

TFNP was first introduced in [MP91].

6.3. Oracle-oriented programming

TODO: unfinished, but this should explain why this chapter is important; another style of programming

If you use reductions + a single oracle for a difficult problem, you can get a powerful programming language. This would be a new, nice paradigm that I aimed to realize.

Due to possible quadratic blowup of output size (even for ac0 circuits i think? but maybe not if fo-uniform?), it is unsatisfactory for us to have a single complete problem solving e.g. sat in worst-time 2^n. Because then it gets 2^{n^2} etc., which is bad. Transductions (or fine-grained time complexity classes such as dlintime) are a potential nice class for this.

6.3.1. Oracle Turing machines and the technique of forcing

TODO: optional. Baker-Gill-Solovay proof uses forcing.

6.3.2. Fine-grained reductions

TODO: just mention QL (quasilinear-time functions), NLT (robust complexity class for $\mathsf{DTIME}(n(\log n)^{\mathcal{O}(1)})$ on RAM machines [GS89])

Chapter 7

Implicit Computational Complexity: recursion-theoretic approach

Implicit computational complexity (ICC) studies how to guarantee resource bounds without appealing to external machine models. Instead of analysing running time or space after the fact, ICC designs languages and recursion schemes whose syntactic constraints ensure that every definable function belongs to a chosen complexity class. The aim is a principled foundation for programming languages that "build in" complexity guarantees by construction. In this chapter we focus on techniques from recursion theory that were successfully utilized in this field to characterize complexity classes.

As discussed in **??**, we will primarily focus on characterizations of L and P. Characterizations we are interested in are described mostly in **??**. All of the historical context of the field needs to also be addressed to describe why we were unsuccessful in designing a programming language based on ideas from these papers. For a broad literature survey, see [Blo94].

7.1. Origins of recursion theory

While not the primary focus of this work, the field of recursion theory developed concepts that later became foundational for ICC. An important formal system studied there is *primitive recursion.*

Definition 7.1.1 (Primitive recursive functions)**.** PR is the smallest class of functions containing (i)–(iii) and closed under (iv) and (v).

(i) (**constants**) for every $n \in \mathbb{N}$ and $k \geq 0$, the k-ary constant function $c_n^{(k)}(\vec{x}) = n$;

(ii) (**successor**) $S(x) = x + 1$;

(iii) (**projections**) for $k \geq 1$ and $1 \leq i \leq k$, $\pi_i^{(k)}(x_1, \ldots, x_k) = x_i$;

(iv) (**composition**) if $h : \mathbb{N}^m \to \mathbb{N}$ and $g_1, \ldots, g_m : \mathbb{N}^k \to \mathbb{N}$ are in PR, then $f(\vec{x}) = h(g_1(\vec{x}), \ldots, g_m(\vec{x}))$ is in PR;

(v) (**primitive recursion**) if $g : \mathbb{N}^k \to \mathbb{N}$ and $h : \mathbb{N}^{k+2} \to \mathbb{N}$ are in PR, then the unique $f : \mathbb{N}^{k+1} \to \mathbb{N}$ is in PR with:

$$f(0, \vec{x}) = g(\vec{x}), \qquad f(S(y), \vec{x}) = h(y, f(y, \vec{x}), \vec{x}).$$

Example 7.1 (Addition)**.** Define $\text{Add} : \mathbb{N}^2 \to \mathbb{N}$ by primitive recursion:

$$\text{Add}(0, x) = x, \qquad \text{Add}(S(y), x) = S(\text{Add}(y, x)).$$

Interestingly, in [MR67] it has been shown that the functions definable in a particular imperative programming language are precisely the primitive recursive functions.

Definition 7.1.2 (LOOP language)**.** Let Var $= \{x_0, x_1, x_2, \dots\}$. LOOP programs are generated by the grammar

$$P ::= x_i := 0 \mid x_i := x_i + 1 \mid P\,; P \mid \texttt{LOOP } x_i \texttt{ DO } P \texttt{ END},$$

where $x_i \in \text{Var}$.

We assume standard semantics, with a remark that $\texttt{LOOP } x_i \texttt{ DO } P \texttt{ END}$ repeats P exactly as many times as the value stored in x_i *at loop entry* (changes to x_i inside P do not change the iteration count).

Theorem 7.2 (LOOP captures precisely PR)**.** *TODO: make statement precise, perhaps based on https://www.cs.cornell.edu/courses/cs6110/2012sp/MeyerAndRitchie-67.pdf*

This simple connection actually satisfies our criteria of a 'programming language capturing a complexity class', as the LOOP language captures exactly PR[1], the class of primitive recursive functions. Moreover, we can even stratify the primitive recursive functions into a hierarchy like in [Grz53].

Historically, the origins of primitive recursion can be traced back to [Gra61] and [Ded88], but the class was probably first considered as the primary object of study in [Sko23]. For the details of the historical origins, consult [Ada11].

7.2. Characterizations not easily adjustable for a programming language

Before the seminal works that founded the field of Implicit Complexity, many characterizations of complexity classes had been known already. All of them suffered at least one of the two problems: either it only characterized a class of relations in a given complexity (as opposed to functions), or the characterization wasn't purely syntactic. We will refer to the latter of being "explicit" instead of "implicit".

7.2.1. Characterizations of classes of relations

Characterizations of classes of relations, such as P (as opposed to FP), are not of interest to us because they don't generalize at all to a programming language allowing to write functions with output. Nevertheless, we investigated the concepts used there and describe some of them briefly in this subsection.

Polynomial-time relations have been characterized without explicit size bounds in [Imm87]. In [CL90], uniform NC[1] was characterized, and in [All91] uniform NC, though their definitions still concealed polynomial bounds and targeted relations instead of functions.

In more modern works, decisive complexity classes have been successfully characterized in [Jon99] by a fragment of Lisp in L and P. The same concept has been extended to account for for nondeterminism in [Bon06].

The authors of [KV05] investigated both imperative and functional programming languages whose fragments yield hierarchies containing *decisional* L, LINSPACE, P, and PSPACE. Related contributions include [Kri05] and [Oit10].

---

[1]As recognized in `https://complexityzoo.net/Complexity_Zoo:P`.

7.2.2. Explicit characterizations

If a characterization of a complexity class is not purely syntactic, i.e. it needs a proof of some additional property besides the function description, it becomes practically impossible to use it as foundation of a practical programming language — it becomes undecidable [2] if a given function description is in the programming language at all. Nevertheless, these concepts have been very important for the field and we shall discuss some of them in this subsection.

An example of such characterization is the Cobham's famous characterization of polynomial-time functions, using the recursion scheme defined below in the style of [Tou22].

Definition 7.2.1. A function f is defined from functions g, h_0, h_1, and s by *bounded primitive recursion on binary notation* if, for every $\vec{x}$ and $y \in \mathbb{N}$,

$$f(\vec{x}, 0) = g(\vec{x}),$$
$$f(\vec{x}, S_0(y)) = h_0(\vec{x}, y, f(\vec{x}, y)),$$
$$f(\vec{x}, S_1(y)) = h_1(\vec{x}, y, f(\vec{x}, y)),$$

and moreover $f(\vec{x}, y) \leq s(\vec{x}, y)$. Here $S_0(y) = 2y$ and $S_1(y) = 2y + 1$ append a binary digit to y.

Definition 7.2.2 (Cobham's algebra for FP)**.** The class of polynomial-time computable functions is the smallest class $\mathcal{C}$ of functions such that:
  (i) the class $\mathcal{C}$ contains the initial functions (as in Definition 7.1.1: constants, successor, projections), the binary successor functions $S_0(x) = 2x$ and $S_1(x) = 2x + 1$, and the weak exponential $(x, y) \mapsto x^{|y|}$;
  (ii) the class $\mathcal{C}$ is closed under composition;
  (iii) the class $\mathcal{C}$ is closed under bounded primitive recursion on binary notation as in Definition Definition 7.2.1.

This formulation, due to Cobham [Cob64], hides the polynomial bounds that are explicit in other presentations. The recursion parameter is the binary representation of the argument, so a definition of $f(\vec{x}, y)$ unfolds only $\mathcal{O}(|y|)$ many steps, which is polynomial in the input size. Moreover, the side condition $f(\vec{x}, y) \leq s(\vec{x}, y)$ forces every intermediate value to stay within numbers of the form $2^{p(|\vec{x}|, |y|)}$ for some polynomial p, hence their binary length remains polynomially bounded. When writing a function in this style, it is unclear how to certify that $f(\vec{x}, y) \leq s(\vec{x}, y)$ holds in a way other than providing a full mathematical proof on the side.

Cobham has published this algebra in [Cob64], suggesting that it captures FP. A proof of that is in [Odi99, p. 175 / p. 186 of the PDF] and in [Tou22, p. 608 / p. 625 of the PDF]

An important characterization is also with function algebras for decisional L and P in Other explicit characterizations include the algebraic view of polynomial-time [Gur83]. Also, as mentioned above in Subsection 7.2.2, uniform NC^1 from [CL90], and uniform NC from [All91].

7.3. Implicit Computational Complexity

We can refine the connection given by Theorem 7.2. As we will see, by carefully modifying the schemes of recursion and composition, we can obtain characterizations of complexity classes such as L and P.

---

[2]Nowadays, proof assistants such as Rocq and Lean could be used to verify a user-provided proof, but this goes out of scope of a *programming language.*

The modern study of ICC begins with two breakthroughs: [Lei91] and [BC92] gave the first implicit characterisations of polynomial-time computable functions. Since then, numerous classes have been captured implicitly; see, for example, [NW10] and [Bon+16] for overviews of FPTIME and FNC characterisations. However, the idea of Bellantoni and Cook seemed to best align with being the foundation of a practical programming language. Hence, we decided to solely focus on it and its successors.

Accessible introductions to ICC include the three-part presentation [Mar06a; Mar06b; Mar06c], the talk [Roc19], and a short overview [Dal12].

7.3.1. Bellantoni and Cook's algebra for FP

Bellantoni and Cook introduced a function algebra $\mathcal{B}$ whose key innovation is the separation of arguments into *normal* inputs (controlling recursion depth) and *safe* inputs (being passed around without influencing that depth). We write $f(\vec{x}; \vec{a})$, with normal inputs $\vec{x}$ to the left of the semicolon and safe inputs $\vec{a}$ to the right. The computation is performed on non-negative integers; proofs transfer to general binary strings (e.g. starting with a zero) [BC92]. For an integer x, let $|x|$ denote its binary length $\lceil \log_2(x+1) \rceil$; for vectors use component-wise notation.

> TODO: dont define |x|, it already appeared in previous subsection

Definition 7.3.1 (Bellantoni-Cook algebra). The class $\mathcal{B}$ is the smallest class of functions on non-negative integers that contains (i)–(v) and is closed under (vi)–(vii).

(i) **(constant)** $0(;) = 0$;

(ii) **(projection)** for $m, n \geq 0$ and $1 \leq j \leq m + n$,

$$\pi_j(x_1, \ldots, x_m; a_1, \ldots, a_n) = \begin{cases} x_j & \text{if } j \leq m, \\ a_{j-m} & \text{otherwise.} \end{cases}$$

;

(iii) **(successors)** $s_i(; a) = 2a + i$ for $i \in \{0, 1\}$;

(iv) **(predecessor)** $p(; 0) = 0$ and $p(; ai) = a$;

(v) **(conditional)**

$$C(; a, b, c) = \begin{cases} b & \text{if } a \bmod 2 = 0, \\ c & \text{otherwise.} \end{cases}$$

;

(vi) **(predicative recursion on notation)** if $g, h_0, h_1 \in \mathcal{B}$, define the new f by

$$f(0, \vec{x}; \vec{a}) = g(\vec{x}; \vec{a}), \qquad f(yi, \vec{x}; \vec{a}) = h_i(y, \vec{x}; \vec{a}, f(y, \vec{x}; \vec{a})),$$

where $i \in \{0, 1\}$ and yi is y with bit i appended.

(vii) **(safe composition)** if $h, \vec{r}, \vec{t} \in \mathcal{B}$ with each component of $\vec{r}$ taking only normal arguments and each component of $\vec{t}$ taking both normal and safe arguments, put

$$f(\vec{x}; \vec{a}) = h(\vec{r}(\vec{x};); \ \vec{t}(\vec{x}; \vec{a})).$$

Intuition. Data from safe arguments, never flow as function output into a normal position when the functions are composed.

Functions in $\mathcal{B}$ can perform arbitrary polynomial-time computation on their normal inputs. Safe inputs may increase only by an additive constant, and recursive results remain safe, so recursion depth cannot depend on previously computed values.

Formalization of polynomial time functions   Interestingly, in [HN11] the authors claim formalizing a proof that a version of Bellantoni and Cook's algebra on bitstrings (as opposed to natural numbers here) captures precisely FP. This is of interest to us, as it might suggest that the authors have formalized the notion of an FP function in a proof assistant and managed to implement some form of an interpreter. This is not the case, however, as their proof is formalized proof that the class of functions in this algebra is the same as the class of functions in Cobham's algebra, as defined in Definition 7.2.2 [3]

7.3.2. Neergaard's safe affine algebra for FL

Møller-Neergaard refines the Bellantoni-Cook discipline to capture FL while insisting that safe data are used *affinely*: each safe value may be consumed at most once. We continue to write functions as $f(\vec{x}; \vec{y})$, with normal inputs $\vec{x}$ and safe inputs $\vec{y}$. Whenever $y \in \mathbb{N}$ and $b \in \{0, 1\}$, write $yb = 2y + b$; for $\delta \in \mathbb{N}$, put $\texttt{shift}(y, \delta) = \lfloor y/2^\delta \rfloor$.

Definition 7.3.2 (Neergaard's BC_ε^- algebra)**.** The class BC_ε^- is the smallest class of functions that contains the initial functions (i)–(v) as defined in Definition 7.3.1 and is closed under the following schemes of safe composition and safe recursion:

(i) **(safe affine composition)** let $(h_i)_{i=1}^M$, $(g_i)_{i=1}^N$, f be functions in BC_ε^- such that each of g_i has only normal arguments ($g_i : \mathbb{N}_2^{m,0}$), and the arity of f is M normal arguments and N safe, i.e. $f : \mathbb{N}_2^{M,N} \to \mathbb{N}_2$. Let also n be the sum of safe arities of h_i.
Define safe affine composition of the functions as a following function in $\mathbb{N}_2^{m,n} \to \mathbb{N}_2$:

$$(f \circ \langle g_1, \ldots, g_M \, ; h_1, \ldots, h_N \rangle)(\vec{x}; \vec{y}) = f(g_1(\vec{x};), \ldots, g_M(\vec{x};); h_1(\vec{x}; \vec{y}_1), \ldots, h_N(\vec{x}; \vec{y}_N)),$$

where $(\vec{y}_1, \ldots, \vec{y}_N)$ partitions the safe inputs into disjoint subtuples; the concatenation $\sum \vec{y}_i$ is equal to $\vec{y}$. ;

<aside>TODO: fix this definition because now its like BC</aside>

(ii) **(safe affine course-of-value recursion)** let $g, h_0, h_1, d_0, d_1 \in BC_\varepsilon^-$. Their course-of-value recursion produces f given by

$$f(0, \vec{x}; \vec{y}) = g(\vec{x}; \vec{y}),$$
$$f(yb, \vec{x}; \vec{y}) = h_b(y, \vec{x}; \vec{y}, f(\text{shift}(y, \delta_b), \vec{x}; \vec{y})),$$

where $b \in \{0, 1\}$ and $\delta_b = |d_b(y, \vec{x};)|$ depends only on the normal part $(y, \vec{x})$.

Intuition.   The recursive result is passed back in a safe position and can be used only once. Safe arguments behave linearly: once a safe value is analysed — for instance by a conditional — it must be recomputed instead of duplicated. Safe affine composition enforces this single-use policy, while the course-of-value recursion permits controlled "look back" determined by the offsets d_0, d_1 computed on normal data. These two ingredients — linearity on safe data and variable-step recursion on normal data — precisely mirror the resources available to logarithmic-space machines; dropping the affinity constraint collapses back to the original Bellantoni-Cook algebra for FP.

---

[3]Link to the source code: `https://github.com/davidnowak/bellantonicook`.

Chapter 8

Linear types

8.1. Implicit Computational Complexity: linear logic approach

`https://github.com/uelis/IntML`

TODO: Here, just show that Ugo dal Lago created IntML. Say roughly what linear types are, don't get into details.

Chapter 9

Bounded arithmetic

In mathematics we typically assume some (pretty strong) foundational axioms we rely on to prove theorems. If we choose set theory as the foundation (as we usually do), a debatable concept is whether we should use the axiom of choice or not. More popularly in computer science, we often want to be explicit about using Kőnig's lemma[1] and Ramsey's theorem. If we think of the concept we introduced in Chapter 6 (i.e. writing a program in a language in which calls of the computation-heavy oracle are explicit), we would often ask ourselves the same question — can we write the program without relying on that function, i.e. write the program in a lower complexity? It turns out the similarity is not a coincidence, and to explore this connection further we need to study the theories of *bounded arithmetic.*

We will start by looking at interesting theorems about one such theory, just after introducing the necessary definitions.

9.1. Single-sorted logic and $I\Delta_0$

Definition 9.1.1 ([CN10, Definition III.1.1]). A *theory* over a vocabulary $\mathcal{L}$ is a set $\mathcal{T}$ of $\mathcal{L}$-formulas that is closed under logical consequence and under universal closure.

Note that we have not defined "logical consequence", which refers to a particular *proof system* (also named *proof calculus*). We will not define proof systems for logic in this work. For those interested, we will just mention that all the results discussed in this chapter assume standard Gentzen-style proof calculus for classical logic, LK [CN10, Section II.2.3] for single-sorted logic and LK^2 [CN10, Section IV.4] for two-sorted logic.

Definition 9.1.2 ([CN10, Definition II.2.3]). The vocabulary of arithmetic is

$$\mathcal{L}_A = \langle 0, 1, +, \cdot \; ; \; =, \leq \rangle.$$

Here $0, 1$ are constant symbols; $+$ and $\cdot$ are binary function symbols; and $=$ and $\leq$ are binary predicate symbols. We will implicitly assume the standard interpretation of these symbols as the appropriate functions on natural numbers whenever talking about the semantics of $\mathcal{L}_A$-formulas.

---

[1] note that Kőnig's lemma is a form of countable choice from finite sets.

Definition 9.1.3 ([CN10, Figure 1] Axioms 1-BASIC of Peano arithmetic)**.**

B1. $x + 1 \neq 0$	B5. $x \cdot 0 = 0$
B2. $x + 1 = y + 1 \to x = y$	B6. $x \cdot (y + 1) = (x \cdot y) + x$
B3. $x + 0 = x$	B7. $(x \leq y \wedge y \leq x) \to x = y$
B4. $x + (y + 1) = (x + y) + 1$	B8. $x \leq x + y$

C. $0 + 1 = 1$

Definition 9.1.4 ([CN10, Definition III.1.4] Induction Scheme)**.** Let Φ be a set of formulas. The Φ-*IND axioms* are all formulas of the form

$$(\varphi(0) \ \wedge \ \forall x. \, (\varphi(x) \to \varphi(x + 1))) \ \to \ \forall z. \, \varphi(z), \tag{9.1}$$

where φ ranges over formulas in Φ. Note that $\varphi(x)$ may have free variables other than x.

Definition 9.1.5 ([CN10, Definition III.1.5] Peano Arithmetic)**.** The theory PA has as axioms $B1, \ldots, B8$, together with the Φ-IND axioms, where Φ is the set of all $\mathcal{L}_A$-formulas.

Peano Arithmetic is a powerful theory capable of formalizing the major theorems of number theory. We define subsystems of PA by restricting the induction axioms to certain sets of formulas.

Definition 9.1.6 ([CN10, Definition III.1.7] IOPEN, $I\Delta_0$, $I\Sigma_1$)**.** Let OPEN be the set of *open* (i.e. quantifier-free) formulas, let Δ_0 be the set of *bounded* formulas, and let Σ_1 be the set of formulas of the form $\exists \vec{x}. \, \varphi$, where φ is bounded and $\vec{x}$ is a (possibly empty) tuple of variables.

The theories IOPEN, $I\Delta_0$, and $I\Sigma_1$ are the subsystems of PA obtained by restricting the induction scheme so that Φ is OPEN, Δ_0, and Σ_1, respectively.

Lemma 9.1 ([CN10, Example III.1.8])**.** *The following formulas (and their universal closures) are theorems of* IOPEN*:*

O1. $(x + y) + z = x + (y + z)$	*(Associativity of +)*
O2. $x + y = y + x$	*(Commutativity of +)*
O3. $x \cdot (y + z) = (x \cdot y) + (x \cdot z)$	*(Distributive law)*
O4. $(x \cdot y) \cdot z = x \cdot (y \cdot z)$	*(Associativity of ·)*
O5. $x \cdot y = y \cdot x$	*(Commutativity of ·)*
O6. $x + z = y + z \to x = y$	*(Cancellation for +)*
O7. $0 \leq x$	
O8. $x \leq 0 \to x = 0$	
O9. $x \leq x$	
O10. $x \neq x + 1$	

Lemma 9.2 ([CN10, Example III.1.9])**.** *The following formulas (and their universal closures)*

are theorems of $I\Delta_0$:

$$
\begin{array}{lll}
\text{D1.} & x \neq 0 \to \exists y \leq x.\,(x = y + 1) & \textit{(Predecessor)} \\
\text{D2.} & \exists z.\,(x + z = y \ \lor\ y + z = x) & \\
\text{D3.} & x \leq y \leftrightarrow \exists z.\,(x + z = y) & \\
\text{D4.} & (x \leq y \land y \leq z) \to x \leq z & \textit{(Transitivity)} \\
\text{D5.} & x \leq y \lor y \leq x & \textit{(Total order)} \\
\text{D6.} & x \leq y \leftrightarrow x + z \leq y + z & \\
\text{D7.} & x \leq y \to x \cdot z \leq y \cdot z & \\
\text{D8.} & x \leq y + 1 \leftrightarrow (x \leq y \lor x = y + 1) & \textit{(Discreteness 1)} \\
\text{D9.} & x < y \leftrightarrow x + 1 \leq y & \textit{(Discreteness 2)} \\
\text{D10.} & x \cdot z = y \cdot z \land z \neq 0 \to x = y & \textit{(Cancellation for} \cdot\textit{)} \\
\end{array}
$$

Using the above lemmas as building blocks, we can prove quite a few nontrivial theorems. We will now introduce the core notion of arithmetic — what does it mean to *define* a function *in a theory*.

Definition 9.1.7 ([CN10, Definition III.3.2] Predicates and Functions definable in a Theory)**.** Let $\mathcal{T}$ be a theory with vocabulary $\mathcal{L}$, and let Φ be a set of $\mathcal{L}$-formulas.

(i) a predicate symbol $P(x) \notin \mathcal{L}$ is Φ-*definable in* $\mathcal{T}$ if there exists an $\mathcal{L}$-formula $\varphi(x) \in \Phi$ such that

$$
P(x) \ \leftrightarrow\ \varphi(x). \tag{9.2}
$$

(ii) a function symbol $f(x) \notin \mathcal{L}$ is Φ-*definable in* $\mathcal{T}$ if there exists a formula $\varphi(x, y) \in \Phi$ such that

$$
\mathcal{T} \vdash \forall x.\,\exists! y.\,\varphi(x, y), \tag{9.3}
$$

and moreover

$$
y = f(x) \ \leftrightarrow\ \varphi(x, y). \tag{9.4}
$$

We call (9.2) a *defining axiom* for $P(x)$ and (9.4) a *defining axiom* for $f(x)$. A symbol is *definable in* $\mathcal{T}$ if it is Φ-definable in $\mathcal{T}$ for some Φ.

Definition 9.1.8. We will say that a function is *provably total* in $\mathcal{T}$ iff it is Σ_1-definable in $\mathcal{T}$.

In [CN10, Section III.3] it is argued that: functions $\lfloor x/y \rfloor$, $\lfloor \sqrt{x} \rfloor$, $\max(0, x - y)$, x mod y are definable in $I\Delta_0$; relation $x \mid y$ is definable in $I\Delta_0$, and, interestingly, the relation $\exp(x, y)$ where $\exp(x, y)$ iff $y = 2^x$, is also definable in $I\Delta_0$. We don't introduce the specific logical formula defining the relation $\exp(x, y)$, as it is complicated and discussed in [CN10, Section III.3]. For a different point of view to these problems, e.g. in [Jum95] it is shown that Euler's φ function is provably total in $I\Delta_0$. However, the limits of expressive power of $I\Delta_0$ are low.

Theorem 9.3 ([CN10, Section III.2])**.**

$$
I\Delta_0 \nvdash \forall x\, \exists y.\, \exp(x, y).
$$

Note that PA *easily proves* $\forall x.\,\exists y.\,\exp(x, y)$.

It is interesting to study the theory $I\Delta_0 + \exp$ of $I\Delta_0$ axioms with an additional axiom stating that the exponential function is definable. As it turns out, this theory enables us

to reason about syntactic constructs such as coding of sets and sequences or context-free grammar parsing [HP93, Chapter V, Section 3] [2].

It turns out that a function is Σ_1-definable in $I\Delta_0$ iff it is in FLTH, functional version of linear-time hierarchy [CN10, Theorem III.4.8]; for the definition of LTH, refer to [CN10, Section III.4.1] — as this complexity class is far from what we call "feasible" in this work, we don't introduce the details here. Instead, we will now introduce a theory with a good computational complexity characterization.

9.2. Two-sorted logic and V^0

Definition 9.2.1 (Axioms of 2-BASIC)**.**

B1. $x + 1 \neq 0$	B7. $(x \leq y \wedge y \leq x) \rightarrow x = y$
B2. $x + 1 = y + 1 \rightarrow x = y$	B8. $x \leq x + y$
B3. $x + 0 = x$	B9. $0 \leq x$
B4. $x + (y + 1) = (x + y) + 1$	B10. $x \leq y \vee y \leq x$
B5. $x \cdot 0 = 0$	B11. $x \leq y \leftrightarrow x < y + 1$
B6. $x \cdot (y + 1) = (x \cdot y) + x$	B12. $x \neq 0 \rightarrow \exists y \leq x.\,(y + 1 = x)$
L1. $X(y) \rightarrow y < |X|$	L2. $y + 1 = |X| \rightarrow X(y)$
SE. $\big(|X| = |Y| \wedge \forall i < |X|.\,(X(i) \leftrightarrow Y(i))\big) \rightarrow X = Y$	

Definition 9.2.2 ([CN10, Definition V.1.2] Comprehension Axiom)**.** Let Φ be a set of formulas. The *comprehension axiom scheme for* Φ, denoted Φ-COMP, consists of all formulas of the form

$$\exists X \leq y.\,\forall z < y.\,\big(X(z) \leftrightarrow \varphi(z)\big), \tag{9.5}$$

where $\varphi(z) \in \Phi$ and X does not occur free in $\varphi(z)$. In (9.5), the formula $\varphi(z)$ may have free variables of both sorts in addition to z. We are mainly interested in the cases where Φ is one of the classes Σ_i^B.

Definition 9.2.3 (V^i)**.** For $i \geq 0$, the theory V^i has vocabulary $\mathcal{L}_A^2$ and is axiomatized by 2-BASIC together with Σ_i^B-COMP.

Note that there are no explicit induction axioms for V^i.

Theorem 9.4 ([CN10, Corollary V.1.8])**.** *Induction is provable in* V^i. *Induction for* Δ_0 *formulas is a theorem of* V^0.

Note that this implies that any theorem φ provable in $I\Delta_0$ *is also provable in* V^0.

Theorem 9.5 ([CN10, Theorem V.1.9])**.** *For every formula φ in the vocabulary $\mathcal{L}_A$ of single-sorted arithmetic, if* $V^0 \vdash \varphi$, *then also* $I\Delta_0 \vdash \varphi$. *In other words,* V^0 *is a* conservative extension *of* $I\Delta_0$.

Remark 9.6 ([CN10, Section IV.3] Two-sorted complexity classes). When operating in two-sorted logic, we need to redefine what does it mean for a relation to be in a complexity class. We will think of numerical arguments x_i of a relation $R(\vec{x}, \vec{X})$ to be passed to the deciding Turing machine in unary representation. The string arguments X_i representing finite sets of numbers are passed as follows. For a string argument S define $S(i) = 1$ when $i \in S$, 0

---

[2]note that they use the name $I\Sigma_0 + \Omega_1$ instead of $I\Delta_0 + \exp$ which is the same.

otherwise. Then the representation $\ulcorner S \urcorner$ of S, when the largest member of S is n, is defined as the following concatenation of bits:

$$\ulcorner S \urcorner = S(n)S(n-1)\ldots S(1)S(0)$$

If S is empty then $\ulcorner S \urcorner$ is the empty string. Note that $|\ulcorner S \urcorner|$ is the same as our interpretation of S inside of the theory: $|S| = \max(S) + 1$ or 0 if S empty.

We will write $(x)_1$ to denote unary representation of x, i.e. $1^{|x|}$ and $(x)_2$ to denote binary representation. The ultimate input to the Turing machine deciding if $R(\vec{x}, \vec{X})$ for $|\vec{x}| = n, |\vec{X}| = N, |X_i| = N_i$ is:

$$(n)_1 \; 0 \quad (x_1)_1 \; 0 \; (x_2)_1 \; 0 \; \ldots \; 0 \; (x_n)_1 \; 0 \quad (N)_1 \; 0 \quad (N_1)_1 \; 0 \; \ulcorner X_1 \urcorner \; 0 \; \ldots \; 0 \; (N_N)_1 \; 0 \; \ulcorner X_N \urcorner$$

Note that a purely numerical relation $R(x)$ is in two-sorted polynomial time iff it is computed in time $2^{\mathcal{O}(n)}$ for $n = \lfloor (x)_2 \rfloor$. The notion of polynomial-time complexity for relations with only string arguments $R(\vec{X})$ coincides with our standard intuition.

Definition 9.2.4 ([CN10, Definition V.2.1]). A number function f or string function F is (*p-bounded*) iff there exists a polynomial $p(x, y)$ such that, for all inputs x, Y,

$$f(x, Y) \; \leq \; p(x, |Y|) \qquad \text{or} \qquad |F(x, Y)| \; \leq \; p(x, |Y|),$$

respectively.

Definition 9.2.5 ([CN10, Definition V.2.3] Two-sorted functional complexity classes). Let C be a two-sorted complexity class of relations. The corresponding *function class* FC consists of:

(i) all *p*-bounded number functions whose graphs belong to C; and

(ii) all *p*-bounded string functions whose bit graphs belong to C.

Note that the classes FAC^0, FP, FL are defined in a different way to what we have used earlier. However, the difference will not matter in this work.

We don't repeat the definitions of definability in a theory for the two-sorted case [CN10, Definition V.4.1]. Recall the definition of Σ_0^B formulas (Definition 2.2.3).

Theorem 9.7 ([CN10, Corollary V.5.3]). *A function is in FAC^0 iff it is Σ_0^B-definable in V^0.*

Definition 9.2.6. The theory VC for a complexity class C has vocabulary $\mathcal{L}_A^2$ and is axiomatized by the axioms of V^0 and one additional axiom depending on the choice of the class C. The additional axiom can be thought of adding an oracle for a C-complete problem to V^0. We skip the (lengthy) technicalities of [CN10, Definition IX.2.1].

The below theorem is the central result of our interest in this thesis.

Theorem 9.8 ([CN10, Theorem IX.2.3]). *A function is provably total in VC iff it is in FC.*

By adding a single axiom to the theory of V^0, we can obtain arithmetical hierarchies in which the functions that we can define and prove correct are precisely the functions from a given complexity class $\mathsf{C} = \mathsf{FTC}^0, \mathsf{FNC}^1, \mathsf{FL}, \mathsf{FP}$.

This way, we obtain theories with very nice properties. They foster certification of complexity of an algorithm (if the proof of correctness itself is feasible, see Subsection 9.2.1). At the same time, they enable us to prove theorems about the correctness of functions defined. In [Bus+25], the authors formalize the breakthrough result $\mathsf{L} = \mathsf{SL}$ of [Rei08] inside of the

weak theory of bounded arithmetic VL. The complexity of computational content of proofs of the Discrete Jordan Curve Theorem is examined in [NC12]. Expander construction in VNC^1 was conducted in [Bus+20].

Another elegant property of these theories is that the proof of a problem not being solvable in a given complexity is exactly a proof of independence of the axiom (corresponding to the problem) from the theory (corresponding to the complexity class).

Theorem 9.9 ([CN06, Corollary 7.21; CN10, Corollary VII.2.4] Independence of PHP from VAC^0)**.**
$$\mathsf{VAC}^0 \nvdash \mathit{PHP}$$

9.2.1. Complexity of algorithm vs complexity of proof

Even when an algorithm is simple, it seems to not always be trivial to "feasibly" prove that it computes the correct result. In our setting, this results in knowing that a problem can be solved in a complexity class C, but not knowing if the corresponding function can be defined in the theory VC (i.e. proved total and correct). See [CN10, Section IX.7.3; CN06, Section 9G.3] for an open problem whether the breakthrough result that binary integer division is in $\mathsf{DLOGTIME}$-uniform TC^0 [HAM02], means that it can also be proved in the corresponding VTC^0 theory. Note that this problem apparently has been solved (affirmatively) in [Jeř22].

Remark 9.10 (Bibliography). The field of bounded arithmetic was initiated by Samuel Buss in his PhD thesis: [Bus86a], in which the theories S_2^1 were introduced to capture reasoning about the polynomial-time hierarchy PH (not introduced in our thesis). The first theory designed to capture polynomial time reasoning was the equational theory PV (as in: polynomially-verifiable [proofs]) theory from [Coo75]. The two-sorted logic language for capturing complexity classes has been introduced by Zambella in [Zam96]. Despite the theories being designed to reason about computation, they are theories of classical logic, which might come off as worrying given our considerations from

TODO: sec: chapter icc, section intuit. logic

. Intuitionistic counterparts such as IS_2^1 for S_2^1 and IPV for PV have also been studied. However, much less is known about their expressive power. For the relation of IS_2^1 and S_2^1, please see [Bus86b]. In particular, [Bus86b, Conjecture 3] asks: if $\mathrm{IS}_2^1 \vdash \exists y. \phi(y, c)$, then is it true that there is a function f, provably correct in S_2^1, such that f *computes* the Gödel encoding of that IS_2^1 proof? In [CU93, Corollary 8.19], that conjecture is answered affirmatively. The intuitionistic version IPV of the theory PV is discussed in some detail in [CU93].

For a good introduction to *bounded reverse mathematics*, with a very thorough overview of arithmetical theories corresponding to complexity classes below FP, refer to [Ngu08].

9.3. Programming language

TODO: In progress

Now, we want to formalize these arithmetical theories so that a computer can check our programs and ensure we didn't go out of a given complexity at any point. With such a programming language (and a formalization of arithmetic in general), we will be able to readily transfer a huge amount of results from paper to computer.

We have two goals for formalization:
 (i) for logicians to believe us the formalization is sound;
 (ii) to be able to extract code with certified complexity from proofs.

There is very little work available on the formalization of arithmetic. A formalization of consistency of Peano arithmetic in Coq was presented in [OCo05]. A formalization of the so-called *Hydra battles* related to unprovability in Peano arithmetic was shown in [Cas+22]. There is an impressive ongoing project of formalization of bounded arithmetic in the model-theoretical style in the Lean community.[3] Their approach doesn't align with our goal of certifying complexity, as their focus is on other arithmetical theories, which differ significantly from what we need. Somehow related, some work on intuitionistic logic in Lean has also been done in [Tru24] even though Lean is not a natural environment for intuitionistic thinking, as it assumes classical axioms very deeply in its standard libraries, unlike Rocq which is constructive by heart.

An idea for a programming language based on bounded arithmetic was discussed in [Li25]. The language they discuss is IMP (PV), based on the equational theory PV which is different from (and less interesting than) the theories we have discussed. There, the authors show how to design an imperative programming language with Hoare logic as the verification mechanism (a.k.a. a type system). Note that for their concept to be implementable in practice, a *formalization* of PV is necessary.

> TODO: Say how i use curry-howard correspondence to extract code from proofs as introduced in chapter linear logic

---

[3] https://github.com/FormalizedFormalLogic/Foundation

```
-- D1.  x ≠ 0 → ∃ y ≤ x, x = y + 1   (Predecessor)
-- proof: induction on x
theorem pred_exists :
  ∀ {x : M}, x ≠ 0 → ∃ y ≤ x, x = y + 1 :=
by
  let ind1 : peano.Formula (Vars2 .y .x) := x =' (y + 1)
  let ind2 : peano.Formula (Vars1 .x) :=
    (Formula.iBdEx' x (display2 .y ind1).flip)
  let ind := idelta0.delta0_induction $ display1 $ (x ≠' 0) ⟹ ind2

  unfold ind2 ind1 at ind

  specialize ind (by
    rw [IsDelta0.display1]
    -- TODO: this lemma can't be in @[delta0_simps],
    -- as it creates a goal 'φ.IsOpen' - which might be not true!
    rw [IsDelta0.of_open.imp]
    · constructor
      · unfold Term.neq
        rw [IsDelta0.of_open.not]
        constructor; constructor; constructor
        constructor; constructor
      · constructor
    · unfold Term.neq
      rw [IsOpen.not]
      constructor; constructor
  )
  simp_induction at ind

  apply ind ?base ?step <;> clear ind ind1 ind2
  · simp only [IsEmpty.forall_iff]
  · intro a hind h
    exists a
    constructor
    · exact B8 a 1
    · rfl
```

Listing 9.1: Lean example

Appendix A

Uniformity

In this chapter, we focus on the descriptions of uniformity conditions used for families of circuits.

A.1. FO-uniformity

We need FO-uniformity for Theorem 5.5. This is defined in [Imm99, Definition 5.16].

A.2. U_{E^*}-uniformity

A notion of U_E and U_{E^*}-uniformity was studied in the early works about uniformity. It is defined in terms of *direct connection language* and *extended connection language* in [Vol99, Definition 2.24, Definition 2.43]. There is a very thorough overview of uniformity conditions below NC^1 in [MIS90]. Since then, this notion has been widely displaced by DLOGTIME-uniformity. Interesting arguments about why DLOGTIME uniformity is the most reasonable to consider are presented in a breakthrogh paper proving that binary integer division is in (DLOGTIME-uniform) TC^0 [HAM02].

A.3. DLOGTIME-uniformity

Source: [MIS90, Section 6]. Important: The class DLOGTIME-uniform NC^1 is equal to ALOGTIME and also equal to is NC^1-uniform NC^1: [MIS90, Lemma 6.2].

There is an example of a concrete DLOGTIME-reduction in [JMT98] (showing that tree isomorphism for string-represented trees is NC^1-complete).

It's mentioned that $complexity AC_0^0 \subseteq$ DLOGTIME $\subseteq \mathsf{AC}_2^0$ [Lee91, Page 141]

Remark A.1 (Bibliography). We refer to a source that is difficult to access (e.g. the above [Lee91, Page 141]): this is Chapter 2 "A Catalog of Complexity Classes" by David S. Johnson [Joh91], which appeared in January 1991 in "Handbook of theoretical computer science (vol. A): algorithms and complexity", edited by Jan van Leeuwen [Lee91].

Bibliography

[AB07]   Sanjeev Arora and Boaz Barak. *Computational Complexity: A Modern Approach*. USA, 2007. URL: https://theory.cs.princeton.edu/complexity/book.pdf.

[AB09]   Sanjeev Arora and Boaz Barak. *Computational Complexity: A Modern Approach*. 1st. USA: Cambridge University Press, 2009. ISBN: 0521424267.

[Ada11]   Rod Adams. *An Early History of Recursive Functions and Computability. From Gödel to Turing*. Reprint/revision of the author's 1983 Ph.D. thesis. Boston: Docent Press, May 28, 2011. 297 pp. ISBN: 978-0-9837004-0-1. URL: https://books.google.com/books?id=G9YoeRIVSnwC (visited on 10/29/2025).

[ADK25]   Melissa Antonelli, Arnaud Durand, and Juha Kontinen. *Characterizing Small Circuit Classes from FAC^0 to FAC^1 via Discrete Ordinary Differential Equations*. 2025. arXiv: 2506.23404 [cs.CC]. URL: https://arxiv.org/abs/2506.23404.

[AKS04]   Manindra Agrawal, Neeraj Kayal, and Nitin Saxena. "PRIMES Is in P". In: *Annals of Mathematics* 160.2 (2004), pp. 781–793. ISSN: 0003486X. URL: http://web.archive.org/web/20250303132722/https://www.cse.iitk.ac.in/users/manindra/algebra/primality_v6.pdf (visited on 11/19/2025).

[All91]   Bill Allen. "Arithmetizing Uniform NC". In: *Annals of Pure and Applied Logic* 53.1 (1991), pp. 1–50. ISSN: 0168-0072. DOI: https://doi.org/10.1016/0168-0072(91)90057-S. URL: https://www.sciencedirect.com/science/article/pii/016800729190057S.

[AR12]   Andrea Asperti and Wilmer Ricciotti. "Formalizing Turing Machines". In: *Logic, Language, Information and Computation*. Ed. by Luke Ong and Ruy de Queiroz. Berlin, Heidelberg: Springer Berlin Heidelberg, 2012, pp. 1–25. ISBN: 978-3-642-32621-9.

[Bal25]   Paweł Balawender. *Seminar Presentation on the Paper "A Formalization of Borel Determinacy in Lean"*. Seminar talk; the paper presented is not by the author. Mar. 2025. URL: https://github.com/ruplet/oracles/blob/f25def5b38a4b1ee60d81c1dcacc765e4fb53595/seminar-presentation-borel-determinacy/prezentacja.pdf (visited on 11/08/2025).

[Bau06]   Andrej Bauer. "First Steps in Synthetic Computability Theory". In: *Electronic Notes in Theoretical Computer Science* 155 (2006). Proceedings of the 21st Annual Conference on Mathematical Foundations of Programming Semantics (MFPS XXI), pp. 5–31. ISSN: 1571-0661. DOI: https://doi.org/10.1016/j.entcs.2005.11.049. URL: https://www.sciencedirect.com/science/article/pii/S1571066106001861.

[BC92]   Stephen Bellantoni and Stephen Cook. "A new recursion-theoretic characterization of the polytime functions". In: *Comput. Complex.* 2.2 (Dec. 1992), pp. 97–110. ISSN: 1016-3328. DOI: 10.1007/BF01201998. URL: https://doi.org/10.1007/BF01201998.

[BD19]   Olivier Bournez and Arnaud Durand. "Recursion Schemes, Discrete Differential Equations and Characterization of Polynomial Time Computations". In: *44th International Symposium on Mathematical Foundations of Computer Science (MFCS 2019)*. Ed. by Peter Rossmanith, Pinar Heggernes, and Joost-Pieter Katoen. Vol. 138. Leibniz International Proceedings in Informatics (LIPIcs). Dagstuhl, Germany: Schloss Dagstuhl – Leibniz-Zentrum für Informatik, 2019, 23:1–23:14. ISBN: 978-3-95977-117-7. DOI: 10.4230/LIPIcs.MFCS.2019.23. URL: https://drops.dagstuhl.de/entities/document/10.4230/LIPIcs.MFCS.2019.23.

[BG92]   M. Bellare and S. Goldwasser. *The Complexity of Decision versus Search*. Tech. rep. USA, 1992. URL: http://web.archive.org/web/20250714180631/https://cseweb.ucsd.edu/~mihir/papers/compip.pdf.

[Blo94]   Stephen Bloch. "Function-algebraic characterizations of log and polylog parallel time". In: *Computational Complexity* 4.2 (June 1994), pp. 175–205. ISSN: 1420-8954. DOI: 10.1007/BF01202288. URL: https://doi.org/10.1007/BF01202288.

[Bob+11]   François Bobot et al. "Why3: Shepherd Your Herd of Provers". In: *Boogie 2011: First International Workshop on Intermediate Verification Languages*. Wrocław, Poland, Aug. 2011, pp. 53–64.

[Boj18]   Mikołaj Bojańczyk. *Polyregular Functions*. 2018. arXiv: 1810.08760 [cs.FL]. URL: https://arxiv.org/abs/1810.08760.

[Bon+16]   Guillaume Bonfante et al. "Two function algebras defining functions in NC^k Boolean circuits". In: *Inf. Comput.* 248.C (June 2016), pp. 82–103. ISSN: 0890-5401. DOI: 10.1016/j.ic.2015.12.009. URL: https://doi.org/10.1016/j.ic.2015.12.009.

[Bon06]   Guillaume Bonfante. "Some Programming Languages for Logspace and Ptime". In: *Algebraic Methodology and Software Technology*. Ed. by Michael Johnson and Varmo Vene. Berlin, Heidelberg: Springer Berlin Heidelberg, 2006, pp. 66–80. ISBN: 978-3-540-35636-3.

[Büc90]   J. Richakd Büchi. "Weak Second-Order Arithmetic and Finite Automata". In: *The Collected Works of J. Richard Büchi*. Ed. by Saunders Mac Lane and Dirk Siefkes. New York, NY: Springer New York, 1990, pp. 398–424. ISBN: 978-1-4613-8928-6. DOI: 10.1007/978-1-4613-8928-6_22. URL: http://web.archive.org/web/20251118183104/https://2024.sci-hub.st/173/04978b7adf46fde189eab42885cdc1d8/10.1002@malq.19600060105.pdf.

[Bus+20]   Sam Buss et al. "Expander construction in VNC1". In: *Annals of Pure and Applied Logic* 171.7 (2020), p. 102796. ISSN: 0168-0072. DOI: https://doi.org/10.1016/j.apal.2020.102796. URL: https://www.sciencedirect.com/science/article/pii/S0168007220300208.

[Bus+25]   Sam Buss et al. *A Logspace Constructive Proof of L=SL*. 2025. arXiv: 2511.12011 [cs.LO]. URL: https://arxiv.org/abs/2511.12011.

[Bus86a]   Samuel R. Buss. *Bounded Arithmetic*. Studies in Proof Theory: Lecture Notes 3. Naples: Bibliopolis, 1986. ISBN: 8870881504. URL: `http://web.archive.org/web/20240421184047/https://mathweb.ucsd.edu/~sbuss/ResearchWeb/BAthesis/Buss_Thesis_OCR.pdf`.

[Bus86b]   Samuel R. Buss. "The polynomial hierarchy and intuitionistic Bounded Arithmetic". In: *Structure in Complexity Theory*. Ed. by Alan L. Selman. Berlin, Heidelberg: Springer Berlin Heidelberg, 1986, pp. 77–103. ISBN: 978-3-540-39825-7.

[Cas+22]   Pierre Castéran et al. "Hydras & Co.: Formalized mathematics in Coq for inspiration and entertainment". In: *Journées Francophones des Langages Applicatifs: JFLA 2022*. St-Médard d'Excideuil, France, June 2022. URL: `https://hal.science/hal-03404668`.

[CFI92]   Jin-Yi Cai, Martin Fürer, and Neil Immerman. "An optimal lower bound on the number of variables for graph identification". In: *Combinatorica* 12.4 (1992), pp. 389–410. ISSN: 1439-6912. DOI: `10.1007/BF01305232`. URL: `https://doi.org/10.1007/BF01305232`.

[Cia16]   Alberto Ciaffaglione. "Towards Turing computability via coinduction". In: *Science of Computer Programming* 126 (2016). Selected Papers from the 17th Brazilian Symposium on Formal Methods (SBMF 2014), pp. 31–51. ISSN: 0167-6423. DOI: `https://doi.org/10.1016/j.scico.2016.02.004`. URL: `https://www.sciencedirect.com/science/article/pii/S0167642316000484`.

[CL90]   Kevin J. Compton and Claude Laflamme. "An algebra and a logic for NC1". In: *Information and Computation* 87.1 (1990). Special Issue: Selections from 1988 IEEE Symposium on Logic in Computer Science, pp. 241–263. ISSN: 0890-5401. DOI: `https://doi.org/10.1016/0890-5401(90)90063-N`. URL: `https://www.sciencedirect.com/science/article/pii/089054019090063N`.

[CM87]   Stephen A Cook and Pierre McKenzie. "Problems complete for deterministic logarithmic space". In: *Journal of Algorithms* 8.3 (1987), pp. 385–394. ISSN: 0196-6774. DOI: `https://doi.org/10.1016/0196-6774(87)90018-6`. eprint: `https://web.archive.org/web/20161213142928/https://www.cs.utoronto.ca/~sacook/homepage/cook_mckenzie.pdf`. URL: `http://web.archive.org/web/20250714180631/https://cseweb.ucsd.edu/~mihir/papers/compip.pdf`.

[CN06]   Stephen A. Cook and Phuong Nguyen. "Logical Foundations of Proof Complexity". Draft manuscript. 2006. URL: `http://web.archive.org/web/20250719042008/https://www.cs.toronto.edu/~sacook/homepage/book/` (visited on 09/03/2025).

[CN10]   Stephen Cook and Phuong Nguyen. *Logical Foundations of Proof Complexity*. Cambridge University Press, 2010. DOI: `10.1017/CBO9780511676277`. URL: `http://web.archive.org/web/20240713034207/https://www.karlin.mff.cuni.cz/~krajicek/cook-nguyen.pdf` (visited on 09/03/2025).

[Cob64]   Alan Cobham. "The Intrinsic Computational Difficulty of Functions". In: *Logic, methodology and philosophy of science*. Ed. by Yehoshua Bar-Hillel. North-Holland Pub. Co., 1964, pp. 24–30. URL: `https://web.archive.org/web/20240121142633/https://www.cs.toronto.edu/~sacook/homepage/cobham_intrinsic.pdf`.

[Coo75]   Stephen A. Cook. "Feasibly constructive proofs and the propositional calculus (Preliminary Version)". In: *Proceedings of the Seventh Annual ACM Symposium on Theory of Computing*. STOC '75. Albuquerque, New Mexico, USA: Association for Computing Machinery, 1975, pp. 83–97. ISBN: 9781450374194. DOI: `10.1145/800116.803756`. URL: `https://doi.org/10.1145/800116.803756`.

[Coo83]   Stephen A. Cook. "An overview of computational complexity". In: *Commun. ACM* 26.6 (June 1983), pp. 400–408. ISSN: 0001-0782. DOI: `10.1145/358141.358144`. URL: `https://doi.org/10.1145/358141.358144`.

[Coo85]   Stephen A. Cook. "A taxonomy of problems with fast parallel algorithms". In: *Information and Control* 64.1 (1985). International Conference on Foundations of Computation Theory, pp. 2–22. ISSN: 0019-9958. DOI: `https://doi.org/10.1016/S0019-9958(85)80041-3`. eprint: `https://web.archive.org/web/20190320200248/https://core.ac.uk/download/pdf/81978561.pdf`. URL: `https://www.sciencedirect.com/science/article/pii/S0019995885800413`.

[Cou94]   Bruno Courcelle. "Monadic second-order definable graph transductions: a survey". In: *Theoretical Computer Science* 126.1 (1994), pp. 53–75. ISSN: 0304-3975. DOI: `https://doi.org/10.1016/0304-3975(94)90268-2`. URL: `https://www.sciencedirect.com/science/article/pii/0304397594902682`.

[CU93]   Stephen Cook and Alasdair Urquhart. "Functional interpretations of feasibly constructive arithmetic". In: *Annals of Pure and Applied Logic* 63.2 (1993), pp. 103–200. ISSN: 0168-0072. DOI: `https://doi.org/10.1016/0168-0072(93)90044-E`. URL: `https://www.sciencedirect.com/science/article/pii/016800729390044E`.

[Dal12]   Ugo Dal Lago. "A Short Introduction to Implicit Computational Complexity". In: *Lectures on Logic and Computation: ESSLLI 2010 Copenhagen, Denmark, August 2010, ESSLLI 2011, Ljubljana, Slovenia, August 2011, Selected Lecture Notes*. Ed. by Nick Bezhanishvili and Valentin Goranko. Berlin, Heidelberg: Springer Berlin Heidelberg, 2012, pp. 89–109. ISBN: 978-3-642-31485-8. DOI: `10.1007/978-3-642-31485-8_3`. URL: `https://doi.org/10.1007/978-3-642-31485-8_3`.

[Dan+01]   Evgeny Dantsin et al. "Complexity and expressive power of logic programming". In: *ACM Comput. Surv.* 33.3 (Sept. 2001), pp. 374–425. ISSN: 0360-0300. DOI: `10.1145/502807.502810`. URL: `http://web.archive.org/web/20240415065931/https://dai.fmph.uniba.sk/~sefranek/bak/dantsin97complexity.pdf`.

[Daw12]   Anuj Dawar. *On Syntactic and Semantic Complexity Classes*. Presentation at the Spitalfields Day, Isaac Newton Institute. Available online via Newton Institute archive. Jan. 2012. URL: `https://web.archive.org/web/20210515020503/https://www.newton.ac.uk/files/seminar/20120109163017301-152985.pdf`.

[Ded88]   Richard Dedekind. *Was sind und was sollen die Zahlen?* German. Erstausgabe 1888; Druckausgabe: Braunschweig, 4. Aufl., 1918. Aachen: semantics Kommunikationsmanagement GmbH, 1888. URL: `https://archive.org/details/WasSindUndWasSollenDieZahlen`.

[DKO21]   Ugo Dal Lago, Reinhard Kahle, and Isabel Oitavem. "A Recursion-Theoretic Characterization of the Probabilistic Class PP". In: *46th International Symposium on Mathematical Foundations of Computer Science (MFCS 2021)*. Ed. by Filippo Bonchi and Simon J. Puglisi. Vol. 202. Leibniz International Proceedings in Informatics (LIPIcs). Dagstuhl, Germany: Schloss Dagstuhl – Leibniz-Zentrum für Informatik, 2021, 35:1–35:12. ISBN: 978-3-95977-201-3. DOI: `10.4230/LIPIcs.MFCS.2021.35`. URL: `https://drops.dagstuhl.de/entities/document/10.4230/LIPIcs.MFCS.2021.35`.

[EH99]   Joost Engelfriet and Hendrik Jan Hoogeboom. *MSO definable string transductions and two-way finite state transducers*. 1999. arXiv: `cs/9906007 [cs.LO]`. URL: `https://arxiv.org/abs/cs/9906007`.

[Elg61]   Calvin C. Elgot. "Decision Problems of Finite Automata Design and Related Arithmetics". In: *Transactions of the American Mathematical Society* 98.1 (1961), pp. 21–51. ISSN: 00029947. URL: `http://web.archive.org/web/20250709105150/https://www.ams.org/journals/tran/1961-098-01/S0002-9947-1961-0139530-9/S0002-9947-1961-0139530-9.pdf` (visited on 11/18/2025).

[Fag74]   Ronald Fagin. "Generalized First-Order Spectra and Polynomial-Time Recognizable Sets". In: *Complexity of Computation*. SIAM-AMS Proceedings 7. American Mathematical Society, 1974, pp. 43–73.

[FKW20]   Yannick Forster, Fabian Kunze, and Maximilian Wuttke. "Verified programming of Turing machines in Coq". In: *Proceedings of the 9th ACM SIGPLAN International Conference on Certified Programs and Proofs*. CPP 2020. New Orleans, LA, USA: Association for Computing Machinery, 2020, pp. 114–128. ISBN: 9781450370974. DOI: `10.1145/3372885.3373816`. URL: `https://doi.org/10.1145/3372885.3373816`.

[For25]   Yannick Forster. "Synthetic Mathematics for the Mechanisation of Computability Theory and Logic". In: *33rd EACSL Annual Conference on Computer Science Logic (CSL 2025)*. Ed. by Jörg Endrullis and Sylvain Schmitz. Vol. 326. Leibniz International Proceedings in Informatics (LIPIcs). Dagstuhl, Germany: Schloss Dagstuhl – Leibniz-Zentrum für Informatik, 2025, 3:1–3:2. ISBN: 978-3-95977-362-1. DOI: `10.4230/LIPIcs.CSL.2025.3`. URL: `https://drops.dagstuhl.de/entities/document/10.4230/LIPIcs.CSL.2025.3`.

[Göb11]   Andreas-Nikolas Göbel. "The Computational Complexity of Computing Nash Equilibrium". Seminar presentation, Algorithmic Game Theory. June 2011. URL: `http://web.archive.org/web/20251109180441/https://old.corelab.ntua.gr/courses/gametheory/old/1011/slides/agob-slides.pdf`.

[Gol21]   Paul W. Goldberg. "Search Problems, and Total Search Problems". Computational Complexity, slides 15; Hilary Term (HT) 2021. 2021. URL: `http://web.archive.org/web/20251109180733/https://www.cs.ox.ac.uk/people/paul.goldberg/CC/2020-21/slides15.pdf`.

[GP18]   Paul W. Goldberg and Christos H. Papadimitriou. "Towards a Unified Complexity Theory of Total Functions". In: *9th Innovations in Theoretical Computer Science Conference, ITCS 2018, January 11-14, 2018, Cambridge, MA, USA*. Ed. by Anna R. Karlin. Vol. 94. LIPIcs. Schloss Dagstuhl - Leibniz-Zentrum für Informatik, 2018, 37:1–37:20. DOI: `10.4230/LIPICS.ITCS.2018.37`. URL: `https://doi.org/10.4230/LIPIcs.ITCS.2018.37`.

[Gra61]   Hermann Grassmann. *Lehrbuch der Arithmetik für höhere Lehranstalten*. German. Verlag von T. C. F. Enslin (Adolph Enslin), 1861, p. 233. URL: https://archive. org/details/lehrbuchderarit00grasgoog/page/n48/mode/2up.

[Grz53]   Andrzej Grzegorczyk. *Some classes of recursive functions*. eng. source: http: //eudml.org/doc/219317. Warszawa: Instytut Matematyczny Polskiej Akademi Nauk, 1953. URL: http://web.archive.org/web/20250806010342/http: //matwbn.icm.edu.pl/ksiazki/rm/rm04/rm0401.pdf.

[GS89]   Yuri Gurevich and Saharon Shelah. "Nearly linear time". In: *Logic at Botik '89*. Ed. by Albert R. Meyer and Michael A. Taitslin. Berlin, Heidelberg: Springer Berlin Heidelberg, 1989, pp. 108–118. ISBN: 978-3-540-46180-7.

[Gur12]   Yuri Gurevich. "What Is an Algorithm?" In: *SOFSEM 2012: Theory and Practice of Computer Science*. Ed. by Mária Bieliková et al. Berlin, Heidelberg: Springer Berlin Heidelberg, 2012, pp. 31–42. ISBN: 978-3-642-27660-6. URL: http:// web.archive.org/web/20240515063009/https://web.eecs.umich.edu/ ~gurevich/Opera/209.pdf.

[Gur83]   Yuri Gurevich. "Algebras of feasible functions". In: *24th Annual Symposium on Foundations of Computer Science (sfcs 1983)*. 1983, pp. 210–214. DOI: 10.1109/ SFCS.1983.5.

[HAM02]   William Hesse, Eric Allender, and David A. Mix Barrington. "Uniform constant-depth threshold circuits for division and iterated multiplication". In: *Journal of Computer and System Sciences* 65.4 (2002). Special Issue on Complexity 2001, pp. 695–716. ISSN: 0022-0000. DOI: https://doi.org/10.1016/S0022-0000(02)00025-9. URL: https://www.sciencedirect.com/science/article/ pii/S0022000002000259.

[Har]   Juris Hartmanis. "Sparse Complete Sets for NP and the Optimal Collapse of the Polynomial Hierarchy". In: *Current Trends in Theoretical Computer Science*, pp. 403–411. DOI: 10.1142/9789812794499_0029. eprint: https://dblp.org/ rec/series/wsscs/Hartmanis93b.html. URL: https://books.google.pl/ books?hl=en&lr=&id=kVhZDTKYa8MC&oi=fnd&pg=PA403&ots=igKOeGCAtC& sig=hnXMtnMszE1gpgVWOa8qaH2-7es&redir_esc=y#v=onepage&q&f=false.

[HK96]   Gerd Hillebrand and Paris Kanellakis. "On the Expressive Power of Simply Typed and Let-Polymorphic Lambda Calculi". In: *Proceedings of the 11th Annual IEEE Symposium on Logic in Computer Science*. LICS '96. USA: IEEE Computer Society, 1996, p. 253. ISBN: 0818674636. URL: http://web.archive.org/web/ 20210506213016/http://citeseerx.ist.psu.edu/viewdoc/download?doi= 10.1.1.23.8845&rep=rep1&type=pdf.

[HKM96]   Gerd G. Hillebrand, Paris C. Kanellakis, and Harry G. Mairson. "Database Query Languages Embedded in the Typed Lambda Calculus". In: *Information and Computation* 127.2 (1996), pp. 117–144. ISSN: 0890-5401. DOI: https://doi.org/ 10.1006/inco.1996.0055. URL: https://www.sciencedirect.com/science/ article/pii/S0890540196900553.

[HN11]   Sylvain Heraud and David Nowak. "A Formalization of Polytime Functions". In: *Interactive Theorem Proving*. Ed. by Marko van Eekelen et al. Berlin, Heidelberg: Springer Berlin Heidelberg, 2011, pp. 119–134. ISBN: 978-3-642-22863-6.

[HP93]   Petr Hájek and Pavel Pudlák. *Metamathematics of First-Order Arithmetic*. 1st ed. Perspectives in Mathematical Logic. Softcover reprint published 17 March 1998. Part of the Springer Book Archive. Berlin, Heidelberg: Springer-Verlag, 1993, pp. XIV+460. ISBN: 978-3-540-63648-9. URL: `https://projecteuclid.org/eBooks/perspectives-in-logic/Metamathematics-of-First-Order-Arithmetic/toc/pl/1235421926`.

[htta]   Auberon (https://cs.stackexchange.com/users/26862/auberon). *Function problems and $FP \subseteq FNP$*. Computer Science Stack Exchange. URL:https://cs.stackexchange.com/q/' (version: 2017-03-16). eprint: `https://cs.stackexchange.com/q/71617`. URL: `https://cs.stackexchange.com/q/71617`.

[httb]   Joshua Grochow (https://cstheory.stackexchange.com/users/129/joshua-grochow). *What exactly are the classes FP, FNP and TFNP?* Theoretical Computer Science Stack Exchange. URL:https://cstheory.stackexchange.com/q/37813 (version: 2017-03-21). eprint: `https://cstheory.stackexchange.com/q/37813`. URL: `https://cstheory.stackexchange.com/q/37813`.

[httc]   Auberon (https://cstheory.stackexchange.com/users/37790/auberon). *What exactly are the classes FP, FNP and TFNP?* Theoretical Computer Science Stack Exchange. URL:https://cstheory.stackexchange.com/q/37812 (version: 2020-06-17). eprint: `https://cstheory.stackexchange.com/q/37812`. URL: `https://cstheory.stackexchange.com/q/37812`.

[httd]   Kaveh (https://mathoverflow.net/users/7507/kaveh). *Is there a syntactic characterization for BPP, BQP, or QMA?* MathOverflow. URL:https://mathoverflow.net/q/35236 (version: 2017-04-13). eprint: `https://mathoverflow.net/q/35236`. URL: `https://mathoverflow.net/q/35236`.

[Imm86]   Neil Immerman. "Relational queries computable in polynomial time". In: *Information and Control* 68.1 (1986), pp. 86–104. ISSN: 0019-9958. DOI: `https://doi.org/10.1016/S0019-9958(86)80029-8`. URL: `https://www.sciencedirect.com/science/article/pii/S0019995886800298`.

[Imm87]   Neil Immerman. "Languages that Capture Complexity Classes". In: *SIAM Journal on Computing* 16.4 (1987), pp. 760–778. DOI: `10.1137/0216051`. eprint: `https://doi.org/10.1137/0216051`. URL: `https://doi.org/10.1137/0216051`.

[Imm99]   Neil Immerman. *Descriptive Complexity*. Springer New York, NY, 1999, p. 268. ISBN: 978-1-4612-0539-5. DOI: `10.1007/978-1-4612-0539-5`. URL: `https://doi.org/10.1007/978-1-4612-0539-5`.

[Jeř22]   Emil Jeřábek. "Iterated Multiplication in VTC^0". In: *Archive for Mathematical Logic* 61.5 (July 2022), pp. 705–767. ISSN: 1432-0665. DOI: `10.1007/s00153-021-00810-6`. URL: `https://doi.org/10.1007/s00153-021-00810-6`.

[JMT98]   B. Jenner, P. McKenzie, and J. Toran. "A note on the hardness of tree isomorphism". In: *Proceedings. Thirteenth Annual IEEE Conference on Computational Complexity (Formerly: Structure in Complexity Theory Conference) (Cat. No.98CB36247)*. 1998, pp. 101–105. DOI: `10.1109/CCC.1998.694595`. eprint: `https://web.archive.org/web/20211202035456/https://www.uni-ulm.de/fileadmin/website_uni_ulm/iui.inst.190/Mitarbeiter/toran/treeiso.pdf`.

[Joh91]   David S. Johnson. "A catalog of complexity classes". In: *Handbook of Theoretical Computer Science (Vol. A): Algorithms and Complexity*. Cambridge, MA, USA: MIT Press, 1991, pp. 67–161. ISBN: 0444880712. URL: http://web.archive.org/web/20231202034313/https://www.dbai.tuwien.ac.at/staff/pichler/complexity/johnson1990.pdf.

[Jon99]   Neil D. Jones. "LOGSPACE and PTIME characterized by programming languages". In: *Theoretical Computer Science* 228.1 (1999), pp. 151–174. ISSN: 0304-3975. DOI: https://doi.org/10.1016/S0304-3975(98)00357-0. URL: https://www.sciencedirect.com/science/article/pii/S0304397598003570.

[Jum95]   Marc Jumelet. "Euler's φ-function in the context of IΔ_0". In: *Archive for Mathematical Logic* 34.3 (June 1995), pp. 197–209. ISSN: 1432-0665. DOI: 10.1007/BF01375521. URL: http://web.archive.org/web/20240910062630/https://dspace.library.uu.nl/bitstream/handle/1874/27697/preprint108.pdf?sequence=1&isAllowed=y.

[Koz06]   Dexter C. Kozen. "The Circuit Value Problem". In: *Theory of Computation*. London: Springer London, 2006, pp. 30–34. ISBN: 978-1-84628-477-9. DOI: 10.1007/1-84628-477-5_6. URL: https://doi.org/10.1007/1-84628-477-5_6.

[Kre88]   Mark W. Krentel. "The complexity of optimization problems". In: *Journal of Computer and System Sciences* 36.3 (1988), pp. 490–509. ISSN: 0022-0000. DOI: https://doi.org/10.1016/0022-0000(88)90039-6. URL: https://www.sciencedirect.com/science/article/pii/0022000088900396.

[Kri05]   Lars Kristiansen. "Neat function algebraic characterizations of logspace and linspace". In: *Computational Complexity* 14.1 (Apr. 2005), pp. 72–88. ISSN: 1420-8954. DOI: 10.1007/s00037-005-0191-0. URL: https://doi.org/10.1007/s00037-005-0191-0.

[Kud96]   Manfred Kudlek. "Small deterministic Turing machines". In: *Theoretical Computer Science* 168.2 (1996), pp. 241–255. ISSN: 0304-3975. DOI: https://doi.org/10.1016/S0304-3975(96)00078-3. URL: https://www.sciencedirect.com/science/article/pii/S0304397596000783.

[Kum+14]   Ramana Kumar et al. "CakeML: A Verified Implementation of ML". In: *Principles of Programming Languages (POPL)*. ACM Press, Jan. 2014, pp. 179–191. DOI: 10.1145/2535838.2535841. URL: https://cakeml.org/popl14.pdf.

[KV05]   Lars Kristiansen and Paul J. Voda. "Programming languages capturing complexity classes". In: *Nordic J. of Computing* 12.2 (Apr. 2005), pp. 89–115. ISSN: 1236-6064. URL: https://www.researchgate.net/publication/220673222_Programming_Languages_Capturing_Complexity_Classes.

[Lad75]   Richard E. Ladner. "The circuit value problem is log space complete for P". In: *SIGACT News* 7.1 (Jan. 1975), pp. 18–20. ISSN: 0163-5700. DOI: 10.1145/990518.990519. URL: https://doi.org/10.1145/990518.990519.

[Lee91]   Jan van Leeuwen, ed. *Handbook of theoretical computer science (vol. A): algorithms and complexity*. Cambridge, MA, USA: MIT Press, 1991. ISBN: 0444880712.

[Lei10]   K. Rustan M. Leino. "Dafny: an automatic program verifier for functional correctness". In: *Proceedings of the 16th International Conference on Logic for Programming, Artificial Intelligence, and Reasoning*. LPAR'10. Dakar, Senegal: Springer-Verlag, 2010, pp. 348–370. ISBN: 3642175104.

[Lei91]   D. Leivant. "A foundational delineation of computational feasibility". In: *[1991] Proceedings Sixth Annual IEEE Symposium on Logic in Computer Science*. 1991, pp. 2–11. DOI: `10.1109/LICS.1991.151625`.

[Ler09]   Xavier Leroy. "A formally verified compiler back-end". In: *Journal of Automated Reasoning* 43.4 (2009), pp. 363–446. URL: `http://xavierleroy.org/publi/compcert-backend.pdf`.

[Li25]   Jiatu Li. *An Introduction to Feasible Mathematics and Bounded Arithmetic for Computer Scientists*. Tech. rep. TR25-086. ISSN 1433-8092. Electronic Colloquium on Computational Complexity (ECCC), Report No. 86, July 2025. URL: `https://eccc.weizmann.ac.il/report/2025/086/`.

[LP99]   Leslie Lamport and Lawrence C. Paulson. "Should your specification language be typed". In: *ACM Trans. Program. Lang. Syst.* 21.3 (May 1999), pp. 502–526. ISSN: 0164-0925. DOI: `10.1145/319301.319317`. URL: `https://doi.org/10.1145/319301.319317`.

[LT12]   Ugo Dal Lago and Paolo Parisen Toldin. *An Higher-Order Characterization of Probabilistic Polynomial Time (Long Version)*. 2012. arXiv: `1202.3317 [cs.LO]`. URL: `https://arxiv.org/abs/1202.3317`.

[Mar06a]   Simone Martini. "Implicit Computational Complexity, part 1". In: *Bertinoro International Spring School for Graduate Studies in Computer Science*. Accessed: 26 August 2025. Bertinoro, Italy, Mar. 2006. URL: `http://web.archive.org/web/20240722203715/https://www.cs.unibo.it/~martini/BISS/martini-1.pdf`.

[Mar06b]   Simone Martini. "Implicit Computational Complexity, part 2". In: *Bertinoro International Spring School for Graduate Studies in Computer Science*. Accessed: 26 August 2025. Bertinoro, Italy, Mar. 2006. URL: `http://web.archive.org/web/20240807183053/http://www.cs.unibo.it/~martini/BISS/martini-2.pdf`.

[Mar06c]   Simone Martini. "Implicit Computational Complexity, part 3". In: *Bertinoro International Spring School for Graduate Studies in Computer Science*. Accessed: 26 August 2025. Bertinoro, Italy, Mar. 2006. URL: `http://web.archive.org/web/20240416100620/http://www.cs.unibo.it/~martini/BISS/martini-3.pdf`.

[Maz18]   Damiano Mazza. *Can typed lambda calculi express *all* algorithms below a given complexity?* Theoretical Computer Science Stack Exchange. URL:https://cstheory.stackexchange.c (version: 2018-04-03). Apr. 3, 2018. eprint: `https://cstheory.stackexchange.com/q/27863`. URL: `https://cstheory.stackexchange.com/q/27863`.

[MIS90]   David A. Mix Barrington, Neil Immerman, and Howard Straubing. "On uniformity within NC1". In: *Journal of Computer and System Sciences* 41.3 (1990), pp. 274–306. ISSN: 0022-0000. DOI: `https://doi.org/10.1016/0022-0000(90)90022-D`. URL: `https://www.sciencedirect.com/science/article/pii/002200009090022D`.

[Mor05]   Tsuyoshi Morioka. "Logical approaches to the complexity of search problems: proof complexity, quantified propositional calculus, and bounded arithmetic". AAINR02726. PhD thesis. CAN, 2005. ISBN: 0494027266. URL: `https://eccc.weizmann.ac.il/resources/pdf/morioka.pdf`.

[MP19]   Anca Muscholl and Gabriele Puppis. "The Many Facets of String Transducers". In: *36th International Symposium on Theoretical Aspects of Computer Science (STACS 2019)*. Ed. by Rolf Niedermeier and Christophe Paul. Vol. 126. Leibniz International Proceedings in Informatics (LIPIcs). Dagstuhl, Germany: Schloss Dagstuhl – Leibniz-Zentrum für Informatik, 2019, 2:1–2:21. ISBN: 978-3-95977-100-9. DOI: 10.4230/LIPIcs.STACS.2019.2. URL: https://drops.dagstuhl.de/entities/document/10.4230/LIPIcs.STACS.2019.2.

[MP91]   Nimrod Megiddo and Christos H. Papadimitriou. "On total functions, existence theorems and computational complexity". In: *Theoretical Computer Science* 81.2 (1991), pp. 317–324. ISSN: 0304-3975. DOI: https://doi.org/10.1016/0304-3975(91)90200-L. URL: https://www.sciencedirect.com/science/article/pii/030439759190200L.

[MR67]   Albert R. Meyer and Dennis M. Ritchie. "The complexity of loop programs". In: *Proceedings of the 1967 22nd National Conference*. ACM '67. Washington, D.C., USA: Association for Computing Machinery, 1967, pp. 465–469. ISBN: 9781450374941. DOI: 10.1145/800196.806014. URL: https://doi.org/10.1145/800196.806014.

[NC12]   Phuong Nguyen and Stephen Cook. "The Complexity of Proving the Discrete Jordan Curve Theorem". In: *ACM Trans. Comput. Logic* 13.1 (Jan. 2012). ISSN: 1529-3785. DOI: 10.1145/2071368.2071377. URL: https://doi.org/10.1145/2071368.2071377.

[Ngu08]   Phuong Nguyen. "Bounded Reverse Mathematics". Electronic Colloquium on Computational Complexity (ECCC) Books, Lectures and Surveys. PhD thesis. Graduate Department of Computer Science, University of Toronto, 2008. URL: https://eccc.weizmann.ac.il/static/books/Bounded_Reverse_Mathematics/.

[NMS21]   Jaroslav Nesetril, Patrice Ossona de Mendez, and Sebastian Siebertz. *Structural properties of the first-order transduction quasiorder*. 2021. arXiv: 2010.02607 [math.CO]. URL: https://arxiv.org/abs/2010.02607.

[NW10]   Karl-Heinz Niggl and Henning Wunderlich. "Implicit characterizations of FPTIME and NC revisited". In: *The Journal of Logic and Algebraic Programming* 79.1 (2010). Special Issue: Logic, Computability and Topology in Computer Science: A New Perspective for Old Disciplines, pp. 47–60. ISSN: 1567-8326. DOI: https://doi.org/10.1016/j.jlap.2009.02.005. URL: https://www.sciencedirect.com/science/article/pii/S1567832609000113.

[OCo05]   Russell O'Connor. "Essential Incompleteness of Arithmetic Verified by Coq". In: *Theorem Proving in Higher Order Logics*. Springer Berlin Heidelberg, 2005, pp. 245–260. ISBN: 9783540318200. DOI: 10.1007/11541868_16. URL: http://dx.doi.org/10.1007/11541868_16.

[Odi99]   Piergiorgio Odifreddi. *Classical Recursion Theory, Volume II*. Vol. 143. Studies in Logic and the Foundations of Mathematics. Amsterdam: Elsevier, 1999, pp. xvi+949. ISBN: 044450205X. URL: https://web.archive.org/web/20240103230629/http://www.piergiorgioodifreddi.it/wp-content/uploads/2010/10/CRT2.pdf.

[Oit10]   Isabel Oitavem. "Logspace without Bounds". In: *Ways of Proof Theory*. Ed. by Ralf Schindler. Berlin, Boston: De Gruyter, 2010, pp. 355–362. ISBN: 9783110324907. DOI: 10.1515/9783110324907.355. URL: https://doi.org/10.1515/9783110324907.355 (visited on 08/27/2025).

[Pap94]   Christos H. Papadimitriou. "On the complexity of the parity argument and other inefficient proofs of existence". In: *Journal of Computer and System Sciences* 48.3 (1994), pp. 498–532. ISSN: 0022-0000. DOI: https://doi.org/10.1016/S0022-0000(05)80063-7. URL: https://www.sciencedirect.com/science/article/pii/S0022000005800637.

[Pau00]   Lawrence C. Paulson. *Set Theory for Verification: II. Induction and Recursion*. 2000. arXiv: cs/9511102 [cs.LO]. URL: https://arxiv.org/abs/cs/9511102.

[Rei08]   Omer Reingold. "Undirected connectivity in log-space". In: *J. ACM* 55.4 (Sept. 2008). ISSN: 0004-5411. DOI: 10.1145/1391289.1391291. URL: https://doi.org/10.1145/1391289.1391291.

[Ric07]   Elaine Rich. *Automata, Computability and Complexity: Theory and Applications*. Prentice-Hall, Sept. 2007. ISBN: 978-0132288064. URL: http://web.archive.org/web/20240407185804/https://www.cs.utexas.edu/~ear/cs341/automatabook/AutomataTheoryBook.pdf.

[Roc19]   Simona Ronchi Della Rocca. "Logic and Implicit Computational Complexity". In: *12th Panhellenic Logic Symposium*. Emerita Professor. Anogeia, Crete, Greece, June 2019. URL: http://web.archive.org/web/20240711151956/http://panhellenic-logic-symposium.org/12/slides/Day1_Ronchi.pdf (visited on 08/26/2025).

[Rog96]   Yurii Rogozhin. "Small universal Turing machines". In: *Theoretical Computer Science* 168.2 (1996), pp. 215–240. ISSN: 0304-3975. DOI: https://doi.org/10.1016/S0304-3975(96)00077-1. URL: https://www.sciencedirect.com/science/article/pii/S0304397596000771.

[Sch05]   Helmut Schmid. "A Programming Language for Finite State Transducers". In: *Finite-State Methods and Natural Language Processing, 5th International Workshop, FSMNLP 2005, Helsinki, Finland, September 1-2, 2005. Revised Papers*. Ed. by Anssi Yli-Jyrä, Lauri Karttunen, and Juhani Karhumäki. Vol. 4002. Lecture Notes in Computer Science. Springer, 2005, pp. 308–309. DOI: 10.1007/11780885_38. URL: https://doi.org/10.1007/11780885_38.

[Sea+25]   Remy Seassau et al. "Formal Semantics and Program Logics for a Fragment of OCaml". In: *Proc. ACM Program. Lang.* 9.ICFP (Aug. 2025). DOI: 10.1145/3747509. URL: https://doi.org/10.1145/3747509.

[Sel94]   Alan L. Selman. "A taxonomy of complexity classes of functions". In: *Journal of Computer and System Sciences* 48.2 (1994), pp. 357–381. ISSN: 0022-0000. DOI: https://doi.org/10.1016/S0022-0000(05)80009-1. URL: https://www.sciencedirect.com/science/article/pii/S0022000005800091.

[Sko23]   Thoralf Skolem. "The foundations of elementary arithmetic established by means of the recursive mode of thought". In: *From Frege to Gödel: A Source Book in Mathematical Logic, 1879–1931*. Ed. by Jean van Heijenoort. English translation from 1967 of Skolem's 1923 paper. Cambridge, MA: Harvard University Press, 1923, pp. 302–333. ISBN: 9780674324497.

[SU06]   Morten Heine Sørensen and Pawel Urzyczyn. *Lectures on the Curry-Howard Isomorphism, Volume 149 (Studies in Logic and the Foundations of Mathematics)*. USA: Elsevier Science Inc., 2006. ISBN: 0444520775.

[Tao10]   Terence Tao. *A Computational Perspective on Set Theory*. Originally published on the blog *What's new*. Archived at http://web.archive.org/web/20250822155936/ https://terrytao.wordpress.com/2010/03/19/a-computational-perspective-on-set-theory/. Mar. 2010. URL: https://terrytao.wordpress.com/2010/03/19/a-computational-perspective-on-set-theory/ (visited on 08/22/2025).

[Tou22]   George Tourlakis. *Computability*. 1st ed. Hardcover ISBN 978-3-030-83201-8 (pub. 2022-08-03); Softcover ISBN 978-3-030-83204-9 (pub. 2023-08-03); eBook ISBN 978-3-030-83202-5 (pub. 2022-08-02); Springer Nature Switzerland AG. Cham: Springer, 2022, pp. XXVII+637. ISBN: 978-3-030-83201-8. DOI: 10.1007/978-3-030-83202-5. URL: https://doi.org/10.1007/978-3-030-83202-5.

[Tru24]   Dafina Trufaş. "Intuitionistic Propositional Logic in Lean". In: *Electronic Proceedings in Theoretical Computer Science* 410 (Oct. 2024), pp. 133–149. ISSN: 2075-2180. DOI: 10.4204/eptcs.410.9. URL: http://dx.doi.org/10.4204/EPTCS.410.9.

[Var82]   Moshe Y. Vardi. "The complexity of relational query languages (Extended Abstract)". In: *Proceedings of the Fourteenth Annual ACM Symposium on Theory of Computing*. STOC '82. San Francisco, California, USA: Association for Computing Machinery, 1982, pp. 137–146. ISBN: 0897910702. DOI: 10.1145/800070.802186. URL: https://doi.org/10.1145/800070.802186.

[Vol99]   Heribert Vollmer. *Introduction to Circuit Complexity: A Uniform Approach*. Berlin, Heidelberg: Springer-Verlag, 1999. ISBN: 3540643109.

[Zak07]   Mateusz Zakrzewski. *Definable functions in the simply typed lambda-calculus*. 2007. arXiv: cs/0701022 [cs.LO]. URL: https://arxiv.org/abs/cs/0701022.

[Zam96]   Domenico Zambella. "Notes on Polynomially Bounded Arithmetic". In: *The Journal of Symbolic Logic* 61.3 (1996), pp. 942–966. ISSN: 00224812. URL: http://www.jstor.org/stable/2275794 (visited on 11/23/2025).